

MIPT & YSDA · AUTUMN 2026

Deep Generative Models

Lecture 4

Roman Isachenko

Moscow Institute of Physics and Technology
Yandex School of Data Analysis

Recap of Previous Lecture

MLE for LVM

$$\sum_{i=1}^n \log p_{\theta}(\mathbf{x}_i) = \sum_{i=1}^n \log \int p_{\theta}(\mathbf{x}_i|\mathbf{z}_i)p(\mathbf{z}_i)d\mathbf{z}_i \rightarrow \max_{\theta}.$$

A Naive Lower Bound

$$\begin{aligned} \log p_{\theta}(\mathbf{x}) &= \log \mathbb{E}_{p(\mathbf{z})} p_{\theta}(\mathbf{x}|\mathbf{z}) \geq \mathbb{E}_{p(\mathbf{z})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) \\ &\approx \frac{1}{K} \sum_{k=1}^K \log p_{\theta}(\mathbf{x}|\mathbf{z}_k), \quad \text{where } \mathbf{z}_k \sim p(\mathbf{z}). \end{aligned}$$

Variational Evidence Lower Bound (ELBO)

$$\begin{aligned} \mathcal{L}_{q,\theta}(\mathbf{x}) &= \mathbb{E}_q \log \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \\ \log p_{\theta}(\mathbf{x}) &= \mathcal{L}_{q,\theta}(\mathbf{x}) + D_{\text{KL}}(q(\mathbf{z})\|p_{\theta}(\mathbf{z}|\mathbf{x})) \geq \mathcal{L}_{q,\theta}(\mathbf{x}). \end{aligned}$$

Recap of Previous Lecture

Variational Evidence Lower Bound (ELBO)

$$\mathcal{L}_{q,\theta}(\mathbf{x}) = \mathbb{E}_q \log \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \leq \log p_{\theta}(\mathbf{x})$$

$$\mathcal{L}_{q,\theta}(\mathbf{x}) = \int q(\mathbf{z}) \log \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} d\mathbf{z} = \mathbb{E}_q \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q(\mathbf{z})\|p(\mathbf{z}))$$

Log-Likelihood Decomposition

$$\log p_{\theta}(\mathbf{x}) = \mathbb{E}_q \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q(\mathbf{z})\|p(\mathbf{z})) + D_{\text{KL}}(q(\mathbf{z})\|p_{\theta}(\mathbf{z}|\mathbf{x})).$$

- Maximize the ELBO as a surrogate for log-likelihood:

$$\max_{\theta} \log p_{\theta}(\mathbf{x}) \longrightarrow \max_{q,\theta} \mathcal{L}_{q,\theta}(\mathbf{x})$$

- Maximizing the ELBO with respect to the variational distribution q is equivalent to minimizing the KL divergence:

$$\arg \max_q \mathcal{L}_{q,\theta}(\mathbf{x}) \equiv \arg \min_q D_{\text{KL}}(q(\mathbf{z})\|p_{\theta}(\mathbf{z}|\mathbf{x})).$$

Recap of Previous Lecture

$$\begin{aligned}\mathcal{L}_{q,\theta}(\mathbf{x}) &= \mathbb{E}_q \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q(\mathbf{z})\|p(\mathbf{z})) \\ &= \mathbb{E}_q \left[\log p_{\theta}(\mathbf{x}|\mathbf{z}) - \log \frac{q(\mathbf{z})}{p(\mathbf{z})} \right] \rightarrow \max_{q,\theta}.\end{aligned}$$

Variational Posterior

$$q^*(\mathbf{z}) = \arg \max_q \mathcal{L}_{q,\theta^*}(\mathbf{x}) = \arg \min_q D_{\text{KL}}(q\|p) = p_{\theta^*}(\mathbf{z}|\mathbf{x}).$$

Amortized Variational Inference

We restrict the family of possible distributions $q(\mathbf{z})$ to a parametric class $q_{\phi}(\mathbf{z}|\mathbf{x})$, conditioned on data $\mathbf{x}$ and parameterized by ϕ .

Gradient Update

$$\begin{bmatrix} \phi_k \\ \theta_k \end{bmatrix} = \begin{bmatrix} \phi_{k-1} + \eta \cdot \nabla_{\phi} \mathcal{L}_{\phi,\theta}(\mathbf{x}) \\ \theta_{k-1} + \eta \cdot \nabla_{\theta} \mathcal{L}_{\phi,\theta}(\mathbf{x}) \end{bmatrix} \Big|_{(\phi_{k-1}, \theta_{k-1})}$$

Recap of Previous Lecture

$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) = \mathbb{E}_q \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z})) \rightarrow \max_{\phi, \theta}.$$

Gradient $\nabla_{\theta} \mathcal{L}_{\phi, \theta}(\mathbf{x})$

$$\begin{aligned} \nabla_{\theta} \mathcal{L}_{\phi, \theta}(\mathbf{x}) &= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \nabla_{\theta} \log p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &\approx \nabla_{\theta} \log p_{\theta}(\mathbf{x}|\mathbf{z}^*), \quad \mathbf{z}^* \sim q_{\phi}(\mathbf{z}|\mathbf{x}). \end{aligned}$$

Gradient $\nabla_{\phi} \mathcal{L}_{\phi, \theta}(\mathbf{x})$

$$\begin{aligned} \nabla_{\phi} \mathcal{L}_{\phi, \theta}(\mathbf{x}) &= \int p(\epsilon) \nabla_{\phi} \log p_{\theta}(\mathbf{x}|\mathbf{g}_{\phi}(\mathbf{x}, \epsilon)) d\epsilon - \nabla_{\phi} D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z})) \\ &\approx \nabla_{\phi} \log p_{\theta}(\mathbf{x}|\mathbf{g}_{\phi}(\mathbf{x}, \epsilon^*)) - \nabla_{\phi} D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z})) \end{aligned}$$

Reparametrization: A Gaussian Example

$$\begin{aligned} p(\epsilon) &= \mathcal{N}(0, \mathbf{I}); \quad q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}_{\phi}(\mathbf{x}), \boldsymbol{\sigma}_{\phi}^2(\mathbf{x})). \\ \mathbf{z} &= \mathbf{g}_{\phi}(\mathbf{x}, \epsilon) = \boldsymbol{\sigma}_{\phi}(\mathbf{x}) \odot \epsilon + \boldsymbol{\mu}_{\phi}(\mathbf{x}). \end{aligned}$$

Recap of Previous Lecture

Training

1. Sample $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$, $\epsilon \sim p(\epsilon)$.
2. Reparametrize $\mathbf{z} = \mathbf{g}_{\phi}(\mathbf{x}, \epsilon)$.
3. Compute the ELBO:

$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) \approx \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})).$$

4. Update ϕ , θ via stochastic gradient ascent.

Sampling

1. Sample $\mathbf{z} \sim p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$.
2. Sample $\mathbf{x} \sim p_{\theta}(\mathbf{x}|\mathbf{z})$.

Note: The encoder $q_{\phi}(\mathbf{z}|\mathbf{x})$ isn't needed during generation.

Outline

01 ELBO Surgery and Optimal VAE Prior

02 Discrete Latent VAEs

Discrete VAE Latent Representations
Vector Quantized VAE (VQ-VAE)

03 Likelihood-Free Learning

04 Generative Adversarial Networks (GAN)

Outline

01 ELBO Surgery and Optimal VAE Prior

02 Discrete Latent VAEs

Discrete VAE Latent Representations
Vector Quantized VAE (VQ-VAE)

03 Likelihood-Free Learning

04 Generative Adversarial Networks (GAN)

ELBO Surgery

$$\frac{1}{n} \sum_{i=1}^n \mathcal{L}_{\phi, \theta}(\mathbf{x}_i) = \frac{1}{n} \sum_{i=1}^n \left[\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}_i)} \log p_{\theta}(\mathbf{x}_i|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z})) \right].$$

Theorem

$$\frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z})) = D_{\text{KL}}(q_{\text{agg}, \phi}(\mathbf{z}) \| p(\mathbf{z})) + \mathbb{I}_q[\mathbf{x}, \mathbf{z}];$$

Aggregated Variational Posterior

$$q_{\text{agg}, \phi}(\mathbf{z}) = \frac{1}{n} \sum_{i=1}^n q_{\phi}(\mathbf{z}|\mathbf{x}_i)$$

- $\mathbb{I}_q[\mathbf{x}, \mathbf{z}]$ is the mutual information between $\mathbf{x}$ and $\mathbf{z}$ under the data distribution $p_{\text{data}}(\mathbf{x})$ and $q_{\phi}(\mathbf{z}|\mathbf{x})$.
- The first term encourages $q_{\text{agg}, \phi}(\mathbf{z})$ to match the prior $p(\mathbf{z})$.
- The second term reduces the information about $\mathbf{x}$ encoded in $\mathbf{z}$.

ELBO Surgery

$$\frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z})) = D_{\text{KL}}(q_{\text{agg},\phi}(\mathbf{z}) \| p(\mathbf{z})) + \mathbb{I}_q[\mathbf{x}, \mathbf{z}].$$

Proof

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z})) &= \frac{1}{n} \sum_{i=1}^n \int q_{\phi}(\mathbf{z}|\mathbf{x}_i) \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x}_i)}{p(\mathbf{z})} d\mathbf{z} \\ &= \frac{1}{n} \sum_{i=1}^n \int q_{\phi}(\mathbf{z}|\mathbf{x}_i) \log \frac{q_{\text{agg},\phi}(\mathbf{z}) q_{\phi}(\mathbf{z}|\mathbf{x}_i)}{p(\mathbf{z}) q_{\text{agg},\phi}(\mathbf{z})} d\mathbf{z} \\ &= \int \frac{1}{n} \sum_{i=1}^n q_{\phi}(\mathbf{z}|\mathbf{x}_i) \log \frac{q_{\text{agg},\phi}(\mathbf{z})}{p(\mathbf{z})} d\mathbf{z} + \frac{1}{n} \sum_{i=1}^n \int q_{\phi}(\mathbf{z}|\mathbf{x}_i) \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x}_i)}{q_{\text{agg},\phi}(\mathbf{z})} d\mathbf{z} \\ &= D_{\text{KL}}(q_{\text{agg},\phi}(\mathbf{z}) \| p(\mathbf{z})) + \frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| q_{\text{agg},\phi}(\mathbf{z})) \\ \mathbb{I}_q[\mathbf{x}, \mathbf{z}] &= \frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| q_{\text{agg},\phi}(\mathbf{z})). \end{aligned}$$

ELBO Surgery

Revisiting the ELBO

$$\begin{aligned}
 \frac{1}{n} \sum_{i=1}^n \mathcal{L}_{\phi, \theta}(\mathbf{x}_i) &= \frac{1}{n} \sum_{i=1}^n \left[\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}_i)} \log p_{\theta}(\mathbf{x}_i|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z})) \right] \\
 &= \underbrace{\frac{1}{n} \sum_{i=1}^n \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}_i)} \log p_{\theta}(\mathbf{x}_i|\mathbf{z})}_{\text{Reconstruction Loss}} - \underbrace{\mathbb{I}_q[\mathbf{x}, \mathbf{z}]}_{\text{Mutual Information}} - \underbrace{D_{\text{KL}}(q_{\text{agg}, \phi}(\mathbf{z}) \| p(\mathbf{z}))}_{\text{Marginal KL}}
 \end{aligned}$$

The **prior distribution** $p(\mathbf{z})$ only appears in the **Marginal KL** term.

Optimal VAE Prior

$$D_{\text{KL}}(q_{\text{agg}, \phi}(\mathbf{z}) \| p(\mathbf{z})) = 0 \quad \Leftrightarrow \quad p(\mathbf{z}) = q_{\text{agg}, \phi}(\mathbf{z}) = \frac{1}{n} \sum_{i=1}^n q_{\phi}(\mathbf{z}|\mathbf{x}_i).$$

Hence, the **optimal prior** $p(\mathbf{z})$ is the **aggregated variational posterior** $q_{\text{agg}, \phi}(\mathbf{z})$.

Marginal KL

$$D_{\text{KL}}(q_{\text{agg},\phi}(\mathbf{z})\|p(\mathbf{z}))$$


Training

The decoder receives $\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x}_i)$; over the training set, $\mathbf{z} \sim q_{\text{agg},\phi}(\mathbf{z})$.

- $q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}_{\phi}(\mathbf{x}), \boldsymbol{\sigma}_{\phi}^2(\mathbf{x}))$ is unimodal.
- A mismatch between the prior and the aggregated posterior can degrade sample quality. Matching these distributions does not by itself eliminate blurriness.

Sampling (generation)

The decoder receives $\mathbf{z} \sim p(\mathbf{z})$, then generates $\mathbf{x} \sim p_{\theta}(\mathbf{x}|\mathbf{z})$.




(a) Prior distribution

(b) Posteriors in standard VAE

Matching the Prior: Which KL Vanishes?


$$\frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z})) = D_{\text{KL}}(q_{\text{agg},\phi}(\mathbf{z}) \| p(\mathbf{z})) + \mathbb{I}_q[\mathbf{x}, \mathbf{z}].$$

$$q_{\phi}(z|\mathbf{x}_i) = \mathcal{N}(\mu_i, 0.65^2), \quad (\mu_1, \mu_2, \mu_3) = (-2, 0, 2), \quad p_{\text{data}}(\mathbf{x}_i) = 1/3$$



Marginal KL	0.584
Mutual information	0.892
Mean conditional KL	1.475

All values in nats; all bars use the same scale.



Marginal KL	0.000
Mutual information	0.892
Mean conditional KL	0.892


All values in nats; all bars use the same scale.

Matching the prior removes **Marginal KL**. The **mutual information** and encoder remain unchanged.

ELBO Surgery in Modern Latent Models

Jointly learn an **autoencoder** and a **generative prior**.

Average conditional KL = **Marginal KL** + **Mutual Information**.



Latents that **reconstruct well** may be **hard for the prior to model**.

GenFirst (2026) studies this balance when the autoencoder and generative prior are trained jointly.

From ELBO to Squared Error

Fix the encoder $q_\phi(\mathbf{z}|\mathbf{x})$, prior $p(\mathbf{z})$ and decoder variance $\sigma^2 > 0$; optimize only θ .

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathcal{L}_{\phi, \theta}(\mathbf{x}) = \mathbb{E}_{p_{\text{data}}(\mathbf{x}) q_\phi(\mathbf{z}|\mathbf{x})} \log p_\theta(\mathbf{x}|\mathbf{z}) - \underbrace{\mathbb{E}_{p_{\text{data}}(\mathbf{x})} D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z}))}_{\text{constant w.r.t. } \theta}.$$

The **conditional KL is constant in θ** , not necessarily zero. At $p = q_{\text{agg}, \phi}$:

$$\underbrace{D_{\text{KL}}(q_{\text{agg}, \phi}(\mathbf{z}) \| p(\mathbf{z}))}_{\text{Marginal KL}} = 0, \quad \underbrace{\frac{1}{n} \sum_{i=1}^n D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}_i) \| p(\mathbf{z}))}_{\text{Average conditional KL}} = \mathbb{I}_q[\mathbf{x}, \mathbf{z}].$$

Gaussian Decoder: Squared Reconstruction Error

$$p_\theta(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\boldsymbol{\mu}_\theta(\mathbf{z}), \sigma^2 \mathbf{I}) \quad \Rightarrow \quad \log p_\theta(\mathbf{x}|\mathbf{z}) = -\frac{\|\mathbf{x} - \boldsymbol{\mu}_\theta(\mathbf{z})\|_2^2}{2\sigma^2} + \text{const.}$$

$$\arg \max_{\theta} \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathcal{L}_{\phi, \theta}(\mathbf{x}) = \arg \min_{\theta} \mathbb{E}_{p_{\text{data}}(\mathbf{x}) q_\phi(\mathbf{z}|\mathbf{x})} \|\mathbf{x} - \boldsymbol{\mu}_\theta(\mathbf{z})\|_2^2.$$

Why are VAE Generations Blurry?

Keep the encoder, prior and Gaussian decoder variance fixed. Maximizing ELBO gives

$$\arg \min_{\theta} \mathbb{E}_{p_{\text{data}}(\mathbf{x})q_{\phi}(\mathbf{z}|\mathbf{x})} \|\mathbf{x} - \boldsymbol{\mu}_{\theta}(\mathbf{z})\|_2^2.$$

Optimal Decoder Mean

The optimal decoder mean is the **conditional expectation**:

$$\boldsymbol{\mu}^*(\mathbf{z}) = \mathbb{E}_{q_{\phi}(\mathbf{x}|\mathbf{z})}[\mathbf{x}] = \frac{\mathbb{E}_{p_{\text{data}}(\mathbf{x})}[q_{\phi}(\mathbf{z}|\mathbf{x}) \cdot \mathbf{x}]}{\mathbb{E}_{p_{\text{data}}(\mathbf{x})}[q_{\phi}(\mathbf{z}|\mathbf{x})]}$$

Blurriness from Averaging

- $p_{\text{data}}(\mathbf{x})$ and $q_{\phi}(\mathbf{z}|\mathbf{x})$ give us the aggregated posterior $q_{\text{agg},\phi}(\mathbf{z})$.
- If $\mathbf{x} \neq \mathbf{x}'$ map to overlapping latent regions, $\boldsymbol{\mu}^*(\mathbf{z})$ averages over unrelated inputs.
- This leads to **blurry, non-distinct outputs**.

Outline

01 ELBO Surgery and Optimal VAE Prior

02 Discrete Latent VAEs

Discrete VAE Latent Representations

Vector Quantized VAE (VQ-VAE)

03 Likelihood-Free Learning

04 Generative Adversarial Networks (GAN)

Discrete VAE Latents

Motivation

- VAE has used **continuous** latent variables $\mathbf{z}$.
- For some modalities, **discrete** representations $\mathbf{z}$ may be a more natural choice.
- Advanced autoregressive models are highly effective for distributions over discrete variables (for example, transformers process discrete tokens).

ELBO

$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z})) \rightarrow \max_{\phi, \theta}.$$

- Apply the reparametrization trick to obtain unbiased gradients.
- Use Gaussian distributions for $q_{\phi}(\mathbf{z}|\mathbf{x})$ and $p(\mathbf{z})$ to compute the KL analytically.

Discrete VAE Latents

Assumptions

- Let $c \sim \text{Cat}(\boldsymbol{\pi})$, where

$$\boldsymbol{\pi} = (\pi_1, \dots, \pi_K), \quad \pi_k = P(c = k), \quad \sum_{k=1}^K \pi_k = 1.$$

- For tokenizer training, fix $p(c) = \text{Uniform}\{1, \dots, K\}$.

ELBO


$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(c|\mathbf{x})} \log p_{\theta}(\mathbf{x}|c) - D_{\text{KL}}(q_{\phi}(c|\mathbf{x}) \| p(c)) \rightarrow \max_{\phi, \theta}.$$

$$\begin{aligned} D_{\text{KL}}(q_{\phi}(c|\mathbf{x}) \| p(c)) &= \sum_{k=1}^K q_{\phi}(k|\mathbf{x}) \log \frac{q_{\phi}(k|\mathbf{x})}{p(k)} \\ &= \sum_{k=1}^K q_{\phi}(k|\mathbf{x}) \log q_{\phi}(k|\mathbf{x}) - \sum_{k=1}^K q_{\phi}(k|\mathbf{x}) \log p(k) \\ &= -H(q_{\phi}(c|\mathbf{x})) + \log K. \end{aligned}$$

Discrete VAE Latents

$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(c|\mathbf{x})} \log p_{\theta}(\mathbf{x}|c) + H(q_{\phi}(c|\mathbf{x})) - \log K \rightarrow \max_{\phi, \theta}.$$

- The encoder should output a discrete distribution $q_{\phi}(c|\mathbf{x})$.
- We need an analogue of the reparametrization trick for discrete $q_{\phi}(c|\mathbf{x})$.
- The decoder $p_{\theta}(\mathbf{x}|c)$ must take a discrete random variable c as input.



Outline

01 ELBO Surgery and Optimal VAE Prior

02 Discrete Latent VAEs

Discrete VAE Latent Representations

Vector Quantized VAE (VQ-VAE)

03 Likelihood-Free Learning

04 Generative Adversarial Networks (GAN)

Vector Quantization

Define the codebook (dictionary) space $\{\mathbf{e}_k\}_{k=1}^K$ with $\mathbf{e}_k \in \mathbb{R}^L$ and K the number of codebook entries.



Quantized Representation

A quantized vector $\mathbf{z}_q \in \mathbb{R}^L$, for any $\mathbf{z} \in \mathbb{R}^L$, is defined via nearest-neighbor lookup in the codebook:

$$\mathbf{z}_q = \mathbf{q}(\mathbf{z}) = \mathbf{e}_{k^*}, \quad \text{where } k^* = \arg \min_k \|\mathbf{z} - \mathbf{e}_k\|.$$

Quantization Procedure

If the encoded tensor has spatial dimensions, quantization is independently applied to each of the $W \times H$ locations.



Vector Quantized VAE (VQ-VAE)

- The encoder outputs a continuous vector $\mathbf{z}_e = \text{NN}_{e,\phi}(\mathbf{x}) \in \mathbb{R}^L$.
- Quantization **deterministically** maps $\mathbf{z}_e$ to its quantized codebook vector $\mathbf{z}_q$.
- The decoder is conditioned on codebook entries $\mathbf{e}_c$, i.e., via $p_\theta(\mathbf{x}|\mathbf{e}_c)$ (instead of $p_\theta(\mathbf{x}|c)$).

Deterministic Variational Posterior

$$q_\phi(c = k^*|\mathbf{x}) = \begin{cases} 1, & \text{for } k^* = \arg \min_k \|\mathbf{z}_e - \mathbf{e}_k\|; \\ 0, & \text{otherwise.} \end{cases}$$

$$D_{\text{KL}}(q_\phi(c|\mathbf{x})\|p(c)) = - \underbrace{\text{H}(q_\phi(c|\mathbf{x}))}_{=0} + \log K = \log K.$$

Constant KL: The aggregated posterior is not pushed toward the uniform prior. We therefore fit a prior to it in a second stage.

Vector Quantized VAE (VQ-VAE): Forward


Deterministic Variational Posterior

$$q_{\phi}(c = k^* | \mathbf{x}) = \begin{cases} 1, & \text{for } k^* = \arg \min_k \|\mathbf{z}_e - \mathbf{e}_k\|; \\ 0, & \text{otherwise.} \end{cases}$$

ELBO

$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(c|\mathbf{x})} \log p_{\theta}(\mathbf{x} | \mathbf{e}_c) - \log K = \log p_{\theta}(\mathbf{x} | \mathbf{z}_q) - \log K,$$

where $\mathbf{z}_q = \mathbf{e}_{k^*}$, $k^* = \arg \min_k \|\mathbf{z}_e - \mathbf{e}_k\|$.




Challenge: The $\arg \min$ operation is non-differentiable.

Vector Quantized VAE (VQ-VAE): Backward

ELBO

$$\mathcal{L}_{\phi, \theta}(\mathbf{x}) = \log p_{\theta}(\mathbf{x}|\mathbf{z}_q) - \log K, \quad \mathbf{z}_q = \mathbf{e}_{k^*}, \quad k^* = \arg \min_k \|\mathbf{z}_e - \mathbf{e}_k\|.$$



Straight-Through Gradient Estimator

$$\begin{aligned} \frac{\partial \log p(\mathbf{x}|\mathbf{z}_q, \theta)}{\partial \phi} &= \frac{\partial \log p_{\theta}(\mathbf{x}|\mathbf{z}_q)}{\partial \mathbf{z}_q} \cdot \frac{\partial \mathbf{z}_q}{\partial \phi} \\ &= \frac{\partial \log p_{\theta}(\mathbf{x}|\mathbf{z}_q)}{\partial \mathbf{z}_q} \cdot \frac{\partial \mathbf{z}_q}{\partial \mathbf{z}_e} \cdot \frac{\partial \mathbf{z}_e}{\partial \phi} \\ &\approx \frac{\partial \log p_{\theta}(\mathbf{x}|\mathbf{z}_q)}{\partial \mathbf{z}_q} \cdot \frac{\partial \mathbf{z}_e}{\partial \phi} \end{aligned}$$

VQ-VAE: Learning a Prior over Codes

The **tokenizer** (encoder + quantizer) produces $M = WH$ codes: $\mathbf{c} = (c_1, \dots, c_M)$.

Target: Aggregated Posterior (ELBO Surgery)

$$p_{\psi}(\mathbf{c}) \approx q_{\text{agg}, \phi}(\mathbf{c}) = \frac{1}{n} \sum_{i=1}^n q_{\phi}(\mathbf{c} | \mathbf{x}_i).$$

Training

1. Fix the trained tokenizer and decoder.
2. Encode the training images $\mathbf{x}_i$ into code maps $\mathbf{c}_i$.
3. Fit the **autoregressive prior** (Lecture 1) by maximum likelihood:

$$\sum_{i=1}^n \log p_{\psi}(\mathbf{c}_i) = \sum_{i=1}^n \sum_{m=1}^M \log p_{\psi}(c_{i,m} | c_{i,1}, \dots, c_{i,m-1}) \rightarrow \max_{\psi}.$$

$p_{\psi}(\mathbf{c}) \rightarrow$ **code map** $\rightarrow$ **codebook lookup** $\rightarrow$ **decoder** $\rightarrow \mathbf{x}$.

Vector Quantized VAE-2 (VQ-VAE-2)

Spatial codes: $\mathbf{c} \in \{1, \dots, K\}^{W \times H}$. VQ-VAE-2 combines code maps at multiple resolutions.

$$q_{\phi}(\mathbf{c}|\mathbf{x}) = \prod_{i=1}^W \prod_{j=1}^H q(c_{ij}|\mathbf{x}, \phi).$$

Sample Diversity: Learned Autoregressive Priors



VQ-VAE (Proposed)

BigGAN deep

Vector Quantized VAE (VQ-VAE): Final algorithm

Training

Stage 1: tokenizer.

1. Sample $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$.
2. Compute encoder outputs $\mathbf{z}_{e,m} = \text{NN}_{e,\phi}(\mathbf{x})_m$.
3. Compute the quantized representation at each location:

$$c_m = \arg \min_k \|\mathbf{z}_{e,m} - \mathbf{e}_k\|_2, \quad \mathbf{z}_{q,m} = \mathbf{e}_{c_m}.$$

4. Compute the ELBO with the uniform training prior:

$$\mathcal{L}_{\phi,\theta}(\mathbf{x}) = \log p_{\theta}(\mathbf{x}|\mathbf{z}_q) - M \log K.$$

5. Compute the loss with codebook and commitment terms:

$$\begin{aligned} \mathcal{L} = & -\mathcal{L}_{\phi,\theta}(\mathbf{x}) + \sum_{m=1}^M \|\text{sg}[\mathbf{z}_{e,m}] - \mathbf{e}_{c_m}\|_2^2 \\ & + \beta \sum_{m=1}^M \|\mathbf{z}_{e,m} - \text{sg}[\mathbf{e}_{c_m}]\|_2^2. \end{aligned}$$

6. Update ϕ , θ and $\{\mathbf{e}_k\}$ (straight-through for ϕ).

Sampling

After **Stage 2: prior training**, fix all parameters.

1. Sample the code map $\mathbf{c} \sim p_{\psi}(\mathbf{c})$ autoregressively.
2. Compute $\mathbf{z}_{q,m} = \mathbf{e}_{c_m}$ at each location.
3. Sample $\mathbf{x} \sim p_{\theta}(\mathbf{x}|\mathbf{z}_q)$.

Outline

01 ELBO Surgery and Optimal VAE Prior

02 Discrete Latent VAEs

Discrete VAE Latent Representations
Vector Quantized VAE (VQ-VAE)

03 **Likelihood-Free Learning**

04 Generative Adversarial Networks (GAN)

Likelihood-Based Models

Poor Likelihood High-Quality Samples

$$p_1(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \mathcal{N}(\mathbf{x} | \mathbf{x}_i, \epsilon \mathbf{I})$$

If ϵ is very small, this model produces excellent, sharp samples but achieves poor likelihoods on test data.

High Likelihood Poor Samples

$$p_2(\mathbf{x}) = 0.01p(\mathbf{x}) + 0.99p_{\text{noise}}(\mathbf{x})$$

$$\log [0.01p(\mathbf{x}) + 0.99p_{\text{noise}}(\mathbf{x})]$$

$$\geq \log [0.01p(\mathbf{x})]$$

$$= \log p(\mathbf{x}) - \log 100$$

This model contains mostly noisy, irrelevant samples; for high dimensions, $\log p(\mathbf{x})$ scales linearly with m .

- Likelihood isn't always a suitable metric for evaluating generative models.
- Sometimes, the likelihood function can't even be computed exactly.

Likelihood-Free Learning

Motivation

We're interested in approximating the true data distribution $p_{\text{data}}(\mathbf{x})$. Instead of searching over all distributions, let's learn a model $p_{\theta}(\mathbf{x}) \approx p_{\text{data}}(\mathbf{x})$.

Suppose we have two sets of samples:

- $\{\mathbf{x}_i\}_{i=1}^{n_1} \sim p_{\text{data}}(\mathbf{x})$ — real data;
- $\{\mathbf{x}_i\}_{i=1}^{n_2} \sim p_{\theta}(\mathbf{x})$ — generated (fake) data.

Define a discriminative model (classifier):

$$p(y = 1|\mathbf{x}) = P(\mathbf{x} \sim p_{\text{data}}(\mathbf{x})); \quad p(y = 0|\mathbf{x}) = P(\mathbf{x} \sim p_{\theta}(\mathbf{x}))$$

Assumption

The generative model $p_{\theta}(\mathbf{x})$ matches $p_{\text{data}}(\mathbf{x})$ if a discriminative model $p(y|\mathbf{x})$ can't distinguish between them — that is, if $p(y = 1|\mathbf{x}) = 0.5$ for every $\mathbf{x}$.

Outline

01 ELBO Surgery and Optimal VAE Prior


02 Discrete Latent VAEs

Discrete VAE Latent Representations
Vector Quantized VAE (VQ-VAE)

03 Likelihood-Free Learning

04 Generative Adversarial Networks (GAN)

Generative Models Taxonomy



Generative Adversarial Networks (GAN)

- The more expressive the discriminator $p(y|\mathbf{x})$, the closer we get to the optimal $p_\theta(\mathbf{x})$.
- Standard classifiers are trained by minimizing cross-entropy loss $-\mathbb{E}_{\hat{p}(\mathbf{x},y)} \log p(y|\mathbf{x})$ with $\hat{p}(\mathbf{x}, y) = \frac{1}{2} \mathbb{I}_{y=1} p_{\text{data}}(\mathbf{x}) + \frac{1}{2} \mathbb{I}_{y=0} p_\theta(\mathbf{x})$.

Cross-Entropy for Discriminator

$$\min_{p(y|\mathbf{x})} \left[-\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log p(y=1|\mathbf{x}) - \mathbb{E}_{p_\theta(\mathbf{x})} \log p(y=0|\mathbf{x}) \right]$$

$$\max_{p(y|\mathbf{x})} \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log p(y=1|\mathbf{x}) + \mathbb{E}_{p_\theta(\mathbf{x})} \log p(y=0|\mathbf{x}) \right]$$

Generative Model

Suppose $p_\theta(\mathbf{x}, \mathbf{z}) = p_\theta(\mathbf{x}|\mathbf{z})p(\mathbf{z})$, where $p(\mathbf{z})$ is a base distribution, and $p_\theta(\mathbf{x}|\mathbf{z}) = \delta(\mathbf{x} - \mathbf{G}_\theta(\mathbf{z}))$ is deterministic.

Connection to VAEs: Both models sample $\mathbf{z} \sim p(\mathbf{z})$. A VAE then samples $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$; here, the decoder is **deterministic**: $\mathbf{x} = \mathbf{G}_\theta(\mathbf{z})$.

Generative Adversarial Networks (GAN)

Cross-Entropy for Discriminative Model

$$\max_{p(y|\mathbf{x})} \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log p(y = 1|\mathbf{x}) + \mathbb{E}_{p_{\theta}(\mathbf{x})} \log p(y = 0|\mathbf{x}) \right]$$

- **Discriminator:** A classifier $p_{\phi}(y = 1|\mathbf{x}) = D_{\phi}(\mathbf{x}) \in [0, 1]$, distinguishing real and generated samples. The discriminator aims to **maximize** cross-entropy.
- **Generator:** The generative model $\mathbf{x} = \mathbf{G}_{\theta}(\mathbf{z})$, $\mathbf{z} \sim p(\mathbf{z})$, seeks to fool the discriminator. The generator aims to **minimize** cross-entropy.

GAN Objective

$$\min_G \max_D \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D(\mathbf{x}) + \mathbb{E}_{p_{\theta}(\mathbf{x})} \log(1 - D(\mathbf{x})) \right]$$

$$\min_G \max_D \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D(\mathbf{G}(\mathbf{z}))) \right]$$

The last step is LOTUS (Law of the Unconscious Statistician): use $\mathbf{x} = \mathbf{G}(\mathbf{z})$ to rewrite the expectation over $p_{\theta}(\mathbf{x})$ as an expectation over $p(\mathbf{z})$.

GAN Optimality

Theorem

The minimax game

$$\min_G \max_D \underbrace{\left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D(\mathbf{G}(\mathbf{z}))) \right]}_{V(G,D)}$$

achieves its global optimum when $p_{\text{data}}(\mathbf{x}) = p_{\theta}(\mathbf{x})$, and $D^*(\mathbf{x}) = 0.5$.

Proof (Fixed G)

$$\begin{aligned} V(G, D) &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D(\mathbf{x}) + \mathbb{E}_{p_{\theta}(\mathbf{x})} \log(1 - D(\mathbf{x})) \\ &= \int \underbrace{[p_{\text{data}}(\mathbf{x}) \log D(\mathbf{x}) + p_{\theta}(\mathbf{x}) \log(1 - D(\mathbf{x}))]}_{y(D)} d\mathbf{x} \\ \frac{dy(D)}{dD} &= \frac{p_{\text{data}}(\mathbf{x})}{D(\mathbf{x})} - \frac{p_{\theta}(\mathbf{x})}{1 - D(\mathbf{x})} = 0 \quad \Rightarrow \quad D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})} \end{aligned}$$

The Optimal Discriminator Responds to the Densities


$$D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_{\theta}(x)}, \quad p_{\text{data}}(x) = \mathcal{N}(0, 1), \quad p_{\theta}(x) = \mathcal{N}(\mu, 1).$$

Separated

 $\mu = 3$

$p_{\text{data}}(x) = \mathcal{N}(0, 1)$

$p_{\theta}(x) = \mathcal{N}(3, 1)$




Equal densities at $x = 1.5$.

Overlap

 $\mu = 1$

$p_{\text{data}}(x) = \mathcal{N}(0, 1)$

$p_{\theta}(x) = \mathcal{N}(1, 1)$




Equal densities at $x = 0.5$.

Matched

 $\mu = 0$

$p_{\text{data}}(x) = \mathcal{N}(0, 1)$

$p_{\theta}(x) = \mathcal{N}(0, 1)$



Equal densities: $D^* = 0.5$ everywhere.

Matching the generator to the data makes the optimal discriminator equal to 0.5 everywhere.

GAN Optimality

Proof Continued (Fixed $D = D^*$)

$$\begin{aligned}
 V(G, D^*) &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log \left(\frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})} \right) + \mathbb{E}_{p_{\theta}(\mathbf{x})} \log \left(\frac{p_{\theta}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})} \right) \\
 &= D_{\text{KL}} \left(p_{\text{data}}(\mathbf{x}) \parallel \frac{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})}{2} \right) + D_{\text{KL}} \left(p_{\theta}(\mathbf{x}) \parallel \frac{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})}{2} \right) - 2 \log 2 \\
 &= 2 D_{\text{JS}}(p_{\text{data}}(\mathbf{x}) \parallel p_{\theta}(\mathbf{x})) - 2 \log 2.
 \end{aligned}$$

Jensen-Shannon Divergence (Symmetric KL Divergence)

$$D_{\text{JS}}(p_{\text{data}}(\mathbf{x}) \parallel p_{\theta}(\mathbf{x})) = \frac{1}{2} \left[D_{\text{KL}} \left(p_{\text{data}}(\mathbf{x}) \parallel \frac{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})}{2} \right) + D_{\text{KL}} \left(p_{\theta}(\mathbf{x}) \parallel \frac{p_{\text{data}}(\mathbf{x}) + p_{\theta}(\mathbf{x})}{2} \right) \right]$$

This can be regarded as a proper distance metric!

$$V(G^*, D^*) = -2 \log 2, \quad p_{\text{data}}(\mathbf{x}) = p_{\theta}(\mathbf{x}), \quad D^*(\mathbf{x}) = 0.5.$$

GAN Optimality

Theorem

The following minimax game

$$\min_G \max_D \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D(\mathbf{G}(\mathbf{z}))) \right]$$

achieves its global optimum precisely when $p_{\text{data}}(\mathbf{x}) = p_{\theta}(\mathbf{x})$, and $D^*(\mathbf{x}) = 0.5$.

Expectations

If the generator can express **any** function and the discriminator is **optimal** at every step, the generator **will converge** to the target distribution.

Reality


- Generator updates are performed in parameter space, and the discriminator is often imperfectly optimized.
- Generator and discriminator losses typically oscillate during GAN training.

GAN Training

Assume both generator and discriminator are parametric models: $D_\phi(\mathbf{x})$ and $\mathbf{G}_\theta(\mathbf{z})$.

Objective

$$\min_{\theta} \max_{\phi} \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D_\phi(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D_\phi(\mathbf{G}_\theta(\mathbf{z}))) \right]$$



- $\mathbf{z} \sim p(\mathbf{z})$ is a latent variable.
- $p_\theta(\mathbf{x}|\mathbf{z}) = \delta(\mathbf{x} - \mathbf{G}_\theta(\mathbf{z}))$ serves as a deterministic decoder (like normalizing flows).
- There is no encoder present.

Summary

- ELBO surgery gives insights into the prior's influence in VAEs; the optimal prior is the aggregated variational posterior.
- With a Gaussian decoder, averaging over distinct inputs that share a latent region can produce blurry outputs even when the prior matches the aggregated posterior.
- VQ-VAE learns discrete codes; a second stage fits their aggregated posterior with a generative prior.
- The straight-through gradient estimator allows gradients to pass as if quantization were an identity operation during backpropagation.
- Likelihood is not always a suitable metric for evaluating generative models; likelihood-free learning motivates the use of discriminative models to match distributions.
- GANs optimize the Jensen-Shannon divergence in theory; training alternates generator and discriminator updates to match the data distribution.