

MIPT & YSDA · AUTUMN 2026

Deep Generative Models

Lecture 6


Roman Isachenko

Moscow Institute of Physics and Technology
Yandex School of Data Analysis

Recap of Previous Lecture

Objective

$$\min_{\theta} \max_{\phi} \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D_{\phi}(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D_{\phi}(\mathbf{G}_{\theta}(\mathbf{z}))) \right]$$



Mode collapse refers to the phenomenon where the generator in a GAN produces only one or a few different modes of the distribution.

Recap of Previous Lecture

Standard GAN Objective

$$\min_{\theta} \max_{\phi} \left[\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \log D_{\phi}(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D_{\phi}(\mathbf{G}_{\theta}(\mathbf{z}))) \right]$$


Theoretical Results

- Both the data distribution $p_{\text{data}}(\mathbf{x})$ and the generative distribution $p_{\theta}(\mathbf{x})$ are supported on low-dimensional manifolds.
- If $p_{\text{data}}(\mathbf{x})$ and $p_{\theta}(\mathbf{x})$ are disjoint, a smooth optimal discriminator can exist!

For such low-dimensional, disjoint manifolds:

$$D_{\text{KL}}(p_{\text{data}} \parallel p_{\theta}) = D_{\text{KL}}(p_{\theta} \parallel p_{\text{data}}) = \infty, \quad D_{\text{JS}}(p_{\text{data}} \parallel p_{\theta}) = \log 2$$

Recap of Previous Lecture



Wasserstein Distance

$$\begin{aligned}
 W(\pi \| p) &= \inf_{\gamma \in \Gamma(\pi, p)} \mathbb{E}_{(\mathbf{x}_1, \mathbf{x}_2) \sim \gamma} \|\mathbf{x}_1 - \mathbf{x}_2\| \\
 &= \inf_{\gamma \in \Gamma(\pi, p)} \int \|\mathbf{x}_1 - \mathbf{x}_2\| \gamma(\mathbf{x}_1, \mathbf{x}_2) d\mathbf{x}_1 d\mathbf{x}_2
 \end{aligned}$$

- $\gamma(\mathbf{x}_1, \mathbf{x}_2)$ is the transport plan: the amount of “dirt” assigned from $\mathbf{x}_1$ to $\mathbf{x}_2$.
- $\Gamma(\pi, p)$ denotes the set of all joint distributions $\gamma(\mathbf{x}_1, \mathbf{x}_2)$ with marginals π and p ($\int \gamma(\mathbf{x}_1, \mathbf{x}_2) d\mathbf{x}_1 = p(\mathbf{x}_2)$, $\int \gamma(\mathbf{x}_1, \mathbf{x}_2) d\mathbf{x}_2 = \pi(\mathbf{x}_1)$).
- $\gamma(\mathbf{x}_1, \mathbf{x}_2)$ is the mass, $\|\mathbf{x}_1 - \mathbf{x}_2\|$ is the distance.

Recap of Previous Lecture

On a compact data space $\mathcal{X}$:

Dependence on Generator Parameters

For a feedforward generator G_θ with Lipschitz activations and $\mathbb{E}_{p(\mathbf{z})} \|\mathbf{z}\| < \infty$, $W(p_{\text{data}} \| p_\theta)$ is continuous in θ and differentiable almost everywhere.

Recap of Previous Lecture

On a compact data space $\mathcal{X}$:

Dependence on Generator Parameters

For a feedforward generator G_θ with Lipschitz activations and $\mathbb{E}_{p(\mathbf{z})} \|\mathbf{z}\| < \infty$, $W(p_{\text{data}} \| p_\theta)$ is continuous in θ and differentiable almost everywhere.

Convergence of Distributions

For a fixed distribution π and a sequence p_t on $\mathcal{X}$, as $t \rightarrow \infty$:

$$D_{\text{KL}}(\pi \| p_t) \rightarrow 0 \implies D_{\text{JS}}(\pi \| p_t) \rightarrow 0 \implies W(\pi \| p_t) \rightarrow 0$$

The first implication also holds for $D_{\text{KL}}(p_t \| \pi) \rightarrow 0$; the converses need not hold.

Wasserstein can capture distributions getting closer even when their supports do not overlap.

Recap of Previous Lecture

Theorem (Kantorovich-Rubinstein Duality)

$$W(p_{\text{data}} \| p_{\theta}) = \frac{1}{K} \max_{\|f\|_L \leq K} [\mathbb{E}_{p_{\text{data}}(\mathbf{x})} f(\mathbf{x}) - \mathbb{E}_{p_{\theta}(\mathbf{x})} f(\mathbf{x})],$$

where $f : \mathbb{R}^m \rightarrow \mathbb{R}$ is K -Lipschitz continuous ($\|f\|_L \leq K$).

WGAN Objective

$$\min_{\theta} W(p_{\text{data}} \| p_{\theta}) \approx \min_{\theta} \max_{\phi \in \Phi} [\mathbb{E}_{p_{\text{data}}(\mathbf{x})} f_{\phi}(\mathbf{x}) - \mathbb{E}_{p(\mathbf{z})} f_{\phi}(\mathbf{G}_{\theta}(\mathbf{z}))].$$

- The discriminator D is replaced by function f : in WGAN, it is known as the **critic**, which is *not* a classifier.
- If the weights ϕ are restricted to a compact set $\Phi = [-c, c]^d$, then $f_{\phi}(\mathbf{x})$ is K -Lipschitz continuous.

$$\begin{aligned} K \cdot W(p_{\text{data}} \| p_{\theta}) &= \max_{\|f\|_L \leq K} [\mathbb{E}_{p_{\text{data}}(\mathbf{x})} f(\mathbf{x}) - \mathbb{E}_{p_{\theta}(\mathbf{x})} f(\mathbf{x})] \\ &\geq \max_{\phi \in \Phi} [\mathbb{E}_{p_{\text{data}}(\mathbf{x})} f_{\phi}(\mathbf{x}) - \mathbb{E}_{p_{\theta}(\mathbf{x})} f_{\phi}(\mathbf{x})] \end{aligned}$$

Recap of Previous Lecture

Frechet Inception Distance (FID)

Use Inception image embeddings $\mathbf{z} = \mathbf{f}(\mathbf{x})$; compare Gaussian approximations of their distributions:

$$\text{FID}(p_{\text{data}}, p_{\theta}) = \|\boldsymbol{\mu}_{\text{data}} - \boldsymbol{\mu}_{\theta}\|^2 + \text{tr} \left[\boldsymbol{\Sigma}_{\text{data}} + \boldsymbol{\Sigma}_{\theta} - 2 \left(\boldsymbol{\Sigma}_{\text{data}}^{1/2} \boldsymbol{\Sigma}_{\theta} \boldsymbol{\Sigma}_{\text{data}}^{1/2} \right)^{1/2} \right].$$

FID depends on the feature network, Gaussian approximation and sample size; it may not correlate with human evaluation.

CMMD: Comparing CLIP Image Embeddings

For CLIP image-embedding distributions P, Q , use a Gaussian kernel k :

$$\text{MMD}_k^2(P, Q) = \mathbb{E}k(\mathbf{z}, \mathbf{z}') + \mathbb{E}k(\mathbf{w}, \mathbf{w}') - 2\mathbb{E}k(\mathbf{z}, \mathbf{w}).$$

Here $\mathbf{z}, \mathbf{z}' \sim P$ and $\mathbf{w}, \mathbf{w}' \sim Q$ are mutually independent. This compares more than means and covariances.


Recap of Previous Lecture

- $\mathcal{S}_{\text{data}} = \{\mathbf{x}_i\}_{i=1}^n \sim p_{\text{data}}(\mathbf{x})$ - real samples;
- $\mathcal{S}_{\theta} = \{\mathbf{x}_i\}_{i=1}^n \sim p_{\theta}(\mathbf{x})$ - generated samples.


Define a binary function:

$$\mathbb{I}(\mathbf{x}, \mathcal{S}) = \begin{cases} 1, & \text{if } \exists \mathbf{x}' \in \mathcal{S} : \|\mathbf{x} - \mathbf{x}'\|_2 \leq \|\mathbf{x}' - \text{NN}_k(\mathbf{x}', \mathcal{S})\|_2; \\ 0, & \text{otherwise.} \end{cases}$$

$$\text{Pr}(\mathcal{S}_{\text{data}}, \mathcal{S}_{\theta}) = \frac{1}{n} \sum_{\mathbf{x} \in \mathcal{S}_{\theta}} \mathbb{I}(\mathbf{x}, \mathcal{S}_{\text{data}}); \quad \text{Rec}(\mathcal{S}_{\text{data}}, \mathcal{S}_{\theta}) = \frac{1}{n} \sum_{\mathbf{x} \in \mathcal{S}_{\text{data}}} \mathbb{I}(\mathbf{x}, \mathcal{S}_{\theta}).$$



(a) True manifold



(b) Approx. manifold

Embed the samples using a pretrained network (as in FID).

Recap of Previous Lecture

Score an image $\mathbf{x}$ for a text prompt t .

CLIP: Learn from Image–Caption Pairs

$$s_{\text{CLIP}}(\mathbf{x}, t) = \cos(\mathbf{f}_{\text{img}}(\mathbf{x}), \mathbf{f}_{\text{text}}(t)) .$$

PickScore: Learn from Human Choices

For the same prompt, a human prefers $\mathbf{x}^+$ to $\mathbf{x}^-$:

$$\mathbb{P}_{\theta}(\mathbf{x}^+ \succ \mathbf{x}^- \mid t) = \frac{\exp r_{\theta}(\mathbf{x}^+, t)}{\exp r_{\theta}(\mathbf{x}^+, t) + \exp r_{\theta}(\mathbf{x}^-, t)} .$$

Train the score r_{θ} to assign higher probability to the human choice.

VQAScore / GenEval / DPG-Bench: evaluate compositional correctness; CLIP similarity can miss counts, attributes and relations.

A learned preference score can guide sample selection and model training — but optimizing the score can exploit its flaws.

Recap of Previous Lecture


- No automated metric is perfect.
- The best way to evaluate generative models is by human assessment.
- It's important to assess various properties.

РЕЛЕВАНТНОСТЬ

«Выполнена ли инструкция?»

Сделал ли редактор именно то, что просили

✓ ХОРОШО



Запрос: «сделать пиратом». Образ пирата создан — то, о чём просили, выполнено.

✗ ПЛОХО




Запрос: эффект вращения на карусели. Картинка почти не изменилась — инструкция не выполнена.

СОХРАНЕНИЕ

«Сохранилось ли то, что менять не просили?»


Фон, лицо и композиция остаются на месте

✓ ХОРОШО



Запрос: образ лётчика, «не меняя черты лица». Под шлемом — то же самое лицо.

✗ ПЛОХО




Запрос: «не меняй внешность, один в один», лишь приблизить лицо к собаке. Но на выходе — совершенно другой человек.

ДЕФЕКТЫ

«Не появилось ли артефактов?»


Лишние пальцы, сломанная геометрия и т. п.

✓ ХОРОШО



Кисти рук и предметы без искажений — артефактов нет.

✗ ПЛОХО



Кисть руки слилась в бесформенный ком (увеличено) — явный артефакт.

Recap of Previous Lecture

Unnormalized Density

$$p_{\theta}(\mathbf{x}) = \frac{\hat{p}_{\theta}(\mathbf{x})}{Z_{\theta}}, \quad \text{where } Z_{\theta} = \int \hat{p}_{\theta}(\mathbf{x}) d\mathbf{x}$$

- $\hat{p}_{\theta}(\mathbf{x})$ can be any non-negative function.
- If we reparameterize as $\hat{p}_{\theta}(\mathbf{x}) = \exp(-f_{\theta}(\mathbf{x}))$, we eliminate the non-negativity constraint.

Recap of Previous Lecture

Unnormalized Density

$$p_{\theta}(\mathbf{x}) = \frac{\hat{p}_{\theta}(\mathbf{x})}{Z_{\theta}}, \quad \text{where } Z_{\theta} = \int \hat{p}_{\theta}(\mathbf{x}) d\mathbf{x}$$

- $\hat{p}_{\theta}(\mathbf{x})$ can be any non-negative function.
- If we reparameterize as $\hat{p}_{\theta}(\mathbf{x}) = \exp(-f_{\theta}(\mathbf{x}))$, we eliminate the non-negativity constraint.

Log-Density Gradient

The gradient of the normalized log-density equals that of the unnormalized log-density:

$$\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) = \nabla_{\mathbf{x}} \log \hat{p}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log Z_{\theta} = \nabla_{\mathbf{x}} \log \hat{p}_{\theta}(\mathbf{x})$$

Recap of Previous Lecture

Unnormalized Density

$$p_{\theta}(\mathbf{x}) = \frac{\hat{p}_{\theta}(\mathbf{x})}{Z_{\theta}}, \quad \text{where } Z_{\theta} = \int \hat{p}_{\theta}(\mathbf{x}) d\mathbf{x}$$

- $\hat{p}_{\theta}(\mathbf{x})$ can be any non-negative function.
- If we reparameterize as $\hat{p}_{\theta}(\mathbf{x}) = \exp(-f_{\theta}(\mathbf{x}))$, we eliminate the non-negativity constraint.

Log-Density Gradient

The gradient of the normalized log-density equals that of the unnormalized log-density:

$$\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) = \nabla_{\mathbf{x}} \log \hat{p}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log Z_{\theta} = \nabla_{\mathbf{x}} \log \hat{p}_{\theta}(\mathbf{x})$$

- Suppose we already have this density (normalized or not) $p_{\theta}(\mathbf{x})$.
- How can we sample from the model?

Recap of Previous Lecture

Theorem (Informal)

Let $\mathbf{x}_0$ be a random vector. Under mild regularity conditions, samples from the following dynamics will eventually follow $p_\theta(\mathbf{x})$ (for sufficiently small η and large l):

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_\theta(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(0, \mathbf{I}).$$


Recap of Previous Lecture

Theorem (Informal)

Let $\mathbf{x}_0$ be a random vector. Under mild regularity conditions, samples from the following dynamics will eventually follow $p_\theta(\mathbf{x})$ (for sufficiently small η and large l):

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_\theta(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(\mathbf{0}, \mathbf{I}).$$

- What if $\boldsymbol{\epsilon}_l = \mathbf{0}$?
- The density $p_\theta(\mathbf{x})$ is the **stationary** distribution of the Markov chain.
- The gradient is taken with respect to $\mathbf{x}$, not $\boldsymbol{\theta}$.
- $\nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})$ defines a vector field.



Outline

01 Score Matching

Denoising Score Matching

Noise-Conditioned Score Network (NCSN)

02 Forward Gaussian Diffusion Process

Outline


01 Score Matching

Denoising Score Matching

Noise-Conditioned Score Network (NCSN)


02 Forward Gaussian Diffusion Process

Generative Models Taxonomy




Score Function

$$\mathbf{s}_\theta(\mathbf{x}) = \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})$$



Score Function


$$\mathbf{s}_\theta(\mathbf{x}) = \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})$$




Langevin Dynamics

$$\begin{aligned} \mathbf{x}_{l+1} &= \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_{\text{data}}(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l \\ &= \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_\theta(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l \\ &= \mathbf{x}_l + \frac{\eta}{2} \cdot \mathbf{s}_\theta(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l. \end{aligned}$$

Score Matching




Score Matching



Fisher Divergence

$$\begin{aligned}
 D_F(p_{\text{data}}, p_\theta) &= \frac{1}{2} \mathbb{E}_{p_{\text{data}}} \left\| \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) \right\|_2^2 \\
 &= \frac{1}{2} \mathbb{E}_{p_{\text{data}}} \left\| \mathbf{s}_\theta(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) \right\|_2^2 \rightarrow \min_{\theta}
 \end{aligned}$$

Score Matching



Fisher Divergence

$$\begin{aligned}
 D_F(p_{\text{data}}, p_\theta) &= \frac{1}{2} \mathbb{E}_{p_{\text{data}}} \left\| \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) \right\|_2^2 \\
 &= \frac{1}{2} \mathbb{E}_{p_{\text{data}}} \left\| \mathbf{s}_\theta(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) \right\|_2^2 \rightarrow \min_{\theta}
 \end{aligned}$$

Problem: We don't know $\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})$.

Why Is the Score Function Better than the Density?

Freedom from Normalization Constants

Many distributions are defined only up to an unnormalized density $\hat{p}(\mathbf{x})$:

$$\nabla_{\mathbf{x}} \log p(\mathbf{x}) = \nabla_{\mathbf{x}} \log \hat{p}(\mathbf{x}) - \underbrace{\nabla_{\mathbf{x}} \log Z}_{=0} = \nabla_{\mathbf{x}} \log \hat{p}(\mathbf{x})$$

The score function bypasses the intractable constant Z entirely!

Why Is the Score Function Better than the Density?

Freedom from Normalization Constants

Many distributions are defined only up to an unnormalized density $\hat{p}(\mathbf{x})$:

$$\nabla_{\mathbf{x}} \log p(\mathbf{x}) = \nabla_{\mathbf{x}} \log \hat{p}(\mathbf{x}) - \underbrace{\nabla_{\mathbf{x}} \log Z}_{=0} = \nabla_{\mathbf{x}} \log \hat{p}(\mathbf{x})$$

The score function bypasses the intractable constant Z entirely!

A Complete Representation

The score function fully characterizes the underlying distribution (up to a constant):

$$\log p(\mathbf{x}) = \log p(\mathbf{x}_0) + \int_0^1 \mathbf{s}(\mathbf{x}_0 + t(\mathbf{x} - \mathbf{x}_0))^{\top} (\mathbf{x} - \mathbf{x}_0) dt$$

Modeling the score is as expressive as modeling $p(\mathbf{x})$ itself, while often more tractable for generative modeling.

Outline

01 Score Matching

Denoising Score Matching

Noise-Conditioned Score Network (NCSN)

02 Forward Gaussian Diffusion Process

Denoising Score Matching

Let us perturb the original data $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$ with Gaussian noise:

$$\mathbf{x}_\sigma = \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}), \quad q(\mathbf{x}_\sigma | \mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I})$$

Denoising Score Matching

Let us perturb the original data $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$ with Gaussian noise:

$$\mathbf{x}_\sigma = \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}), \quad q(\mathbf{x}_\sigma | \mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I})$$

$$q(\mathbf{x}_\sigma) = \int q(\mathbf{x}_\sigma | \mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}.$$

Denoising Score Matching

Let us perturb the original data $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$ with Gaussian noise:

$$\mathbf{x}_\sigma = \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}), \quad q(\mathbf{x}_\sigma | \mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I})$$

$$q(\mathbf{x}_\sigma) = \int q(\mathbf{x}_\sigma | \mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}.$$

Assumption

The solution to

$$\frac{1}{2} \mathbb{E}_{q(\mathbf{x}_\sigma)} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) \right\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

satisfies $\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) \approx \mathbf{s}_{\boldsymbol{\theta}, 0}(\mathbf{x}_0) = \mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x})$ if σ is sufficiently small.

Denoising Score Matching

Let us perturb the original data $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$ with Gaussian noise:

$$\mathbf{x}_\sigma = \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}), \quad q(\mathbf{x}_\sigma | \mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I})$$

$$q(\mathbf{x}_\sigma) = \int q(\mathbf{x}_\sigma | \mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}.$$

Assumption

The solution to

$$\frac{1}{2} \mathbb{E}_{q(\mathbf{x}_\sigma)} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) \right\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

satisfies $\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) \approx \mathbf{s}_{\boldsymbol{\theta}, 0}(\mathbf{x}_0) = \mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x})$ if σ is sufficiently small.

- The score function of the noised data nearly matches the score function of the original data.
- The score function $\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma)$ is parameterized by σ .
- **Note:** We don't know $q(\mathbf{x}_\sigma)$, just as we don't know $p_{\text{data}}(\mathbf{x})$.

Denoising Score Matching

Theorem

Under mild regularity conditions, the following holds:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_\sigma)} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) \right\|_2^2 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma | \mathbf{x})} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}) \right\|_2^2 + \text{const}(\boldsymbol{\theta}) \end{aligned}$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the following holds:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_\sigma)} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) \right\|_2^2 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma | \mathbf{x})} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}) \right\|_2^2 + \text{const}(\boldsymbol{\theta}) \end{aligned}$$

Gradient of the Noise Kernel

$$\mathbf{x}_\sigma = \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_\sigma | \mathbf{x}) = \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I})$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the following holds:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_\sigma)} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) \right\|_2^2 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma | \mathbf{x})} \left\| \mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}) \right\|_2^2 + \text{const}(\boldsymbol{\theta}) \end{aligned}$$

Gradient of the Noise Kernel

$$\begin{aligned} \mathbf{x}_\sigma &= \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, & q(\mathbf{x}_\sigma | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I}) \\ \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}) &= -\frac{\mathbf{x}_\sigma - \mathbf{x}}{\sigma^2} = -\frac{\boldsymbol{\epsilon}}{\sigma} \end{aligned}$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the following holds:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_\sigma)} \left\| \mathbf{s}_{\theta, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) \right\|_2^2 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma | \mathbf{x})} \left\| \mathbf{s}_{\theta, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}) \right\|_2^2 + \text{const}(\boldsymbol{\theta}) \end{aligned}$$

Gradient of the Noise Kernel

$$\begin{aligned} \mathbf{x}_\sigma &= \mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}, & q(\mathbf{x}_\sigma | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma^2 \cdot \mathbf{I}) \\ \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}) &= -\frac{\mathbf{x}_\sigma - \mathbf{x}}{\sigma^2} = -\frac{\boldsymbol{\epsilon}}{\sigma} \end{aligned}$$

- The right-hand side doesn't require computing $\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)$ or even $\nabla_{\mathbf{x}_\sigma} \log p_{\text{data}}(\mathbf{x}_\sigma)$.
- $\mathbf{s}_{\theta, \sigma}(\mathbf{x}_\sigma)$ is trained to **denoise** the noised samples $\mathbf{x}_\sigma$.

Denoising Score Matching

Initial objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Denoising Score Matching

Initial objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Noised objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Denoising Score Matching

Initial objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Noised objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

This is equivalent to a denoising task:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_{\sigma}|\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}_{\sigma}} \log q(\mathbf{x}_{\sigma}|\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Denoising Score Matching

Initial objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Noised objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\theta, \sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\theta}$$

This is equivalent to a denoising task:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_{\sigma}|\mathbf{x})} \|\mathbf{s}_{\theta, \sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}_{\sigma}} \log q(\mathbf{x}_{\sigma}|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{\mathcal{N}(0, \mathbf{I})} \left\| \mathbf{s}_{\theta, \sigma}(\mathbf{x} + \sigma \cdot \boldsymbol{\epsilon}) + \frac{\boldsymbol{\epsilon}}{\sigma} \right\|_2^2 \rightarrow \min_{\theta}$$

Langevin Dynamics

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \frac{\eta}{2} \cdot \mathbf{s}_{\theta, \sigma}(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(0, \mathbf{I}).$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof

For any integrable h , marginalize the joint distribution:

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} h(\mathbf{x}_\sigma) = \int \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) h(\mathbf{x}_\sigma) d\mathbf{x}_\sigma = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} h(\mathbf{x}_\sigma).$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof

For any integrable h , marginalize the joint distribution:

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} h(\mathbf{x}_\sigma) = \int \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) h(\mathbf{x}_\sigma) d\mathbf{x}_\sigma = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} h(\mathbf{x}_\sigma).$$

Expand the squared error:

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof

For any integrable h , marginalize the joint distribution:

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} h(\mathbf{x}_\sigma) = \int \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) h(\mathbf{x}_\sigma) d\mathbf{x}_\sigma = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} h(\mathbf{x}_\sigma).$$

Expand the squared error:

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 \\ &= \mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma)\|_2^2 + \underbrace{\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2}_{C_1} - 2\mathbb{E}_{q(\mathbf{x}_\sigma)} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)]. \end{aligned}$$

The score-norm term C_1 is independent of θ .

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof (Continued)

Rewrite the mixed term:

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)] = \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma) d\mathbf{x}_\sigma$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof (Continued)

Rewrite the mixed term:

$$\begin{aligned} \mathbb{E}_{q(\mathbf{x}_\sigma)} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)] &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma) d\mathbf{x}_\sigma \\ &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) d\mathbf{x}_\sigma = \iint p_{\text{data}}(\mathbf{x}) \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma|\mathbf{x}) d\mathbf{x} d\mathbf{x}_\sigma \end{aligned}$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof (Continued)

Rewrite the mixed term:

$$\begin{aligned} \mathbb{E}_{q(\mathbf{x}_\sigma)} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)] &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma) d\mathbf{x}_\sigma \\ &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) d\mathbf{x}_\sigma = \iint p_{\text{data}}(\mathbf{x}) \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma|\mathbf{x}) d\mathbf{x} d\mathbf{x}_\sigma \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})] . \end{aligned}$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof (Continued)

Rewrite the mixed term:

$$\begin{aligned} \mathbb{E}_{q(\mathbf{x}_\sigma)} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)] &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma) d\mathbf{x}_\sigma \\ &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) d\mathbf{x}_\sigma = \iint p_{\text{data}}(\mathbf{x}) \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma|\mathbf{x}) d\mathbf{x} d\mathbf{x}_\sigma \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]. \end{aligned}$$

Substitute and complete the square:

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} [\|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma)\|_2^2 - 2\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})] + C_1$$

Denoising Score Matching

Theorem

Under mild regularity conditions, the last term is independent of θ :

$$\mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta).$$

Proof (Continued)

Rewrite the mixed term:

$$\begin{aligned} \mathbb{E}_{q(\mathbf{x}_\sigma)} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)] &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma) d\mathbf{x}_\sigma \\ &= \int \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \left(\int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x} \right) d\mathbf{x}_\sigma = \iint p_{\text{data}}(\mathbf{x}) \mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} q(\mathbf{x}_\sigma|\mathbf{x}) d\mathbf{x} d\mathbf{x}_\sigma \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} [\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]. \end{aligned}$$

Substitute and complete the square:

$$\begin{aligned} \mathbb{E}_{q(\mathbf{x}_\sigma)} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)\|_2^2 &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} [\|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma)\|_2^2 - 2\mathbf{s}_{\theta,\sigma}^\top(\mathbf{x}_\sigma) \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})] + C_1 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 + \text{const}(\theta). \end{aligned}$$

Denoising Score Matching

Denoising Score Matching Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma | \mathbf{x})} \|\mathbf{s}_{\theta, \sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Denoising Score Matching

Denoising Score Matching Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

The optimal MSE predictor is the posterior mean of the regression target:

$$\mathbf{s}_{\theta^*,\sigma}(\mathbf{x}_\sigma) = \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]$$

Denoising Score Matching

Denoising Score Matching Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

The optimal MSE predictor is the posterior mean of the regression target:

$$\mathbf{s}_{\theta^*,\sigma}(\mathbf{x}_\sigma) = \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]$$

Marginal Score via Conditional Scores

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) = \frac{\nabla_{\mathbf{x}_\sigma} \int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}}{q(\mathbf{x}_\sigma)}$$

Denoising Score Matching

Denoising Score Matching Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

The optimal MSE predictor is the posterior mean of the regression target:

$$\mathbf{s}_{\theta^*,\sigma}(\mathbf{x}_\sigma) = \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]$$

Marginal Score via Conditional Scores

$$\begin{aligned} \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) &= \frac{\nabla_{\mathbf{x}_\sigma} \int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}}{q(\mathbf{x}_\sigma)} \\ &= \int \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x}) \cdot \frac{q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x})}{q(\mathbf{x}_\sigma)} d\mathbf{x} \end{aligned}$$

Denoising Score Matching

Denoising Score Matching Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

The optimal MSE predictor is the posterior mean of the regression target:

$$\mathbf{s}_{\theta^*,\sigma}(\mathbf{x}_\sigma) = \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]$$

Marginal Score via Conditional Scores

$$\begin{aligned} \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) &= \frac{\nabla_{\mathbf{x}_\sigma} \int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}}{q(\mathbf{x}_\sigma)} \\ &= \int \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x}) \cdot \frac{q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x})}{q(\mathbf{x}_\sigma)} d\mathbf{x} \\ &= \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})] \end{aligned}$$

Denoising Score Matching

Denoising Score Matching Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_\sigma|\mathbf{x})} \|\mathbf{s}_{\theta,\sigma}(\mathbf{x}_\sigma) - \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

The optimal MSE predictor is the posterior mean of the regression target:

$$\mathbf{s}_{\theta^*,\sigma}(\mathbf{x}_\sigma) = \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})]$$

Marginal Score via Conditional Scores

$$\begin{aligned} \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) &= \frac{\nabla_{\mathbf{x}_\sigma} \int q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x}) d\mathbf{x}}{q(\mathbf{x}_\sigma)} \\ &= \int \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x}) \cdot \frac{q(\mathbf{x}_\sigma|\mathbf{x}) p_{\text{data}}(\mathbf{x})}{q(\mathbf{x}_\sigma)} d\mathbf{x} \\ &= \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma|\mathbf{x})] \end{aligned}$$

Therefore: $\mathbf{s}_{\theta^*,\sigma}(\mathbf{x}_\sigma) = \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)$ - the optimal model learns the true **marginal** score!

Tweedie's Formula

Let $\mathbf{x} \sim p_{\text{data}}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$ be independent, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma \boldsymbol{\epsilon}$ and $\sigma > 0$.

Using the posterior-mean identity from the previous slide:

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) = \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} \left[\frac{\mathbf{x} - \mathbf{x}_\sigma}{\sigma^2} \right]$$

Tweedie's Formula

Let $\mathbf{x} \sim p_{\text{data}}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$ be independent, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma \boldsymbol{\epsilon}$ and $\sigma > 0$.

Using the posterior-mean identity from the previous slide:

$$\begin{aligned}\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) &= \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} \left[\frac{\mathbf{x} - \mathbf{x}_\sigma}{\sigma^2} \right] \\ &= \frac{\mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} [\mathbf{x}] - \mathbf{x}_\sigma}{\sigma^2}.\end{aligned}$$

Tweedie's Formula

Let $\mathbf{x} \sim p_{\text{data}}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$ be independent, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma\boldsymbol{\epsilon}$ and $\sigma > 0$.

Using the posterior-mean identity from the previous slide:

$$\begin{aligned}\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) &= \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} \left[\frac{\mathbf{x} - \mathbf{x}_\sigma}{\sigma^2} \right] \\ &= \frac{\mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)}[\mathbf{x}] - \mathbf{x}_\sigma}{\sigma^2}.\end{aligned}$$

Tweedie's Formula

$$\mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)}[\mathbf{x}] = \mathbf{x}_\sigma + \sigma^2 \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma).$$

Tweedie's Formula

Let $\mathbf{x} \sim p_{\text{data}}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$ be independent, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma\boldsymbol{\epsilon}$ and $\sigma > 0$.

Using the posterior-mean identity from the previous slide:

$$\begin{aligned}\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) &= \mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)} \left[\frac{\mathbf{x} - \mathbf{x}_\sigma}{\sigma^2} \right] \\ &= \frac{\mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)}[\mathbf{x}] - \mathbf{x}_\sigma}{\sigma^2}.\end{aligned}$$

Tweedie's Formula

$$\mathbb{E}_{q(\mathbf{x}|\mathbf{x}_\sigma)}[\mathbf{x}] = \mathbf{x}_\sigma + \sigma^2 \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma).$$

Noise Prediction


$$\mathbb{E}[\boldsymbol{\epsilon} \mid \mathbf{x}_\sigma] = -\sigma \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma).$$

Both conditional means are MSE-optimal predictions given the noisy observation.

From Conditional Scores to Denoising

Toy data: three equally likely clean points $\mathbf{x}_i$, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma\epsilon$, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$ and $\sigma = 1$.

Fixed noisy observation between the clean-data modes.




Conditional scores

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}_i) = (\mathbf{x}_i - \mathbf{x}_\sigma) / \sigma^2$$

From Conditional Scores to Denoising

Toy data: three equally likely clean points $\mathbf{x}_i$, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma\epsilon$, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$ and $\sigma = 1$.

Fixed noisy observation between the clean-data modes.



Conditional scores

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}_i) = (\mathbf{x}_i - \mathbf{x}_\sigma) / \sigma^2$$

Posterior weights

$$q(\mathbf{x}_1 | \mathbf{x}_\sigma)$$

22.9%

$$q(\mathbf{x}_2 | \mathbf{x}_\sigma)$$

42.7%


$$q(\mathbf{x}_3 | \mathbf{x}_\sigma)$$

34.4%

From Conditional Scores to Denoising

Toy data: three equally likely clean points $\mathbf{x}_i$, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma\epsilon$, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$ and $\sigma = 1$.


Fixed noisy observation between the clean-data modes.



Conditional scores

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}_i) = (\mathbf{x}_i - \mathbf{x}_\sigma) / \sigma^2$$

Posterior weights




Marginal score = posterior-weighted average

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) = \sum_i q(\mathbf{x}_i | \mathbf{x}_\sigma) \frac{\mathbf{x}_i - \mathbf{x}_\sigma}{\sigma^2}$$

From Conditional Scores to Denoising

Toy data: three equally likely clean points $\mathbf{x}_i$, with $\mathbf{x}_\sigma = \mathbf{x} + \sigma\epsilon$, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$ and $\sigma = 1$.

Fixed noisy observation between the clean-data modes.



Conditional scores

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma | \mathbf{x}_i) = (\mathbf{x}_i - \mathbf{x}_\sigma) / \sigma^2$$

Posterior weights

$$q(\mathbf{x}_1 | \mathbf{x}_\sigma)$$

22.9%

$$q(\mathbf{x}_2 | \mathbf{x}_\sigma)$$

42.7%

$$q(\mathbf{x}_3 | \mathbf{x}_\sigma)$$

34.4%

Marginal score = posterior-weighted average

$$\nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma) = \sum_i q(\mathbf{x}_i | \mathbf{x}_\sigma) \frac{\mathbf{x}_i - \mathbf{x}_\sigma}{\sigma^2}$$

Tweedie: from the score to the clean-data mean

$$\mathbb{E}[\mathbf{x} | \mathbf{x}_\sigma] = \mathbf{x}_\sigma + \sigma^2 \nabla_{\mathbf{x}_\sigma} \log q(\mathbf{x}_\sigma)$$

Posterior mean: (-0.35, 0.26)

The MSE-optimal denoised prediction can lie between the clean-data modes.

Denoising Score Matching

Original objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Denoising Score Matching

Original objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Noisy objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\boldsymbol{\theta}, \sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Denoising Score Matching

Original objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Noisy objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

This is equivalent to a denoising task:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_{\sigma}|\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}_{\sigma}} \log q(\mathbf{x}_{\sigma}|\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Denoising Score Matching

Original objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Noisy objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

This is equivalent to a denoising task:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_{\sigma}|\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}_{\sigma}} \log q(\mathbf{x}_{\sigma}|\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{\mathcal{N}(0,\mathbf{I})} \left\| \mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x} + \sigma \boldsymbol{\epsilon}) + \frac{\boldsymbol{\epsilon}}{\sigma} \right\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Denoising Score Matching

Original objective:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Noisy objective:

$$\mathbb{E}_{q(\mathbf{x}_{\sigma})} \|\mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}} \log q(\mathbf{x}_{\sigma})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

This is equivalent to a denoising task:

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_{\sigma}|\mathbf{x})} \|\mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_{\sigma}) - \nabla_{\mathbf{x}_{\sigma}} \log q(\mathbf{x}_{\sigma}|\mathbf{x})\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{\mathcal{N}(0,\mathbf{I})} \left\| \mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x} + \sigma \boldsymbol{\epsilon}) + \frac{\boldsymbol{\epsilon}}{\sigma} \right\|_2^2 \rightarrow \min_{\boldsymbol{\theta}}$$

Langevin Dynamics

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \frac{\eta}{2} \cdot \mathbf{s}_{\boldsymbol{\theta},\sigma}(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(0, \mathbf{I})$$

Outline

01 Score Matching

Denoising Score Matching

Noise-Conditioned Score Network (NCSN)


02 Forward Gaussian Diffusion Process

Denoising Score Matching

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{\mathcal{N}(0, \mathbf{I})} \left\| \mathbf{s}_{\theta, \sigma}(\mathbf{x} + \sigma \boldsymbol{\epsilon}) + \frac{\boldsymbol{\epsilon}}{\sigma} \right\|_2^2 \rightarrow \min_{\theta}$$

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \frac{\eta}{2} \cdot \mathbf{s}_{\theta, \sigma}(\mathbf{x}_l) + \sqrt{\eta} \cdot \boldsymbol{\epsilon}_l$$

- For **small** σ , $\mathbf{s}_{\theta, \sigma}(\mathbf{x})$ becomes inaccurate and Langevin dynamics fails to traverse modes.
- For **large** σ , robustness in low-density regions is achieved, but the model learns a distribution that is overly corrupted.



Noise-Conditioned Score Network (NCSN)


- Specify a sequence of noise levels: $\sigma_1 < \sigma_2 < \dots < \sigma_T$.
- Perturb each data point with different noise levels: $\mathbf{x}_t = \mathbf{x} + \sigma_t \boldsymbol{\epsilon}$, so $\mathbf{x}_t \sim q(\mathbf{x}_t)$.
- Choose σ_1, σ_T such that:

$$q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \sigma_T^2 \mathbf{I})$$

Noise-Conditioned Score Network (NCSN)

- Specify a sequence of noise levels: $\sigma_1 < \sigma_2 < \dots < \sigma_T$.
- Perturb each data point with different noise levels: $\mathbf{x}_t = \mathbf{x} + \sigma_t \epsilon$, so $\mathbf{x}_t \sim q(\mathbf{x}_t)$.
- Choose σ_1, σ_T such that:

$$q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \sigma_T^2 \mathbf{I})$$



Noise-Conditioned Score Network (NCSN)

Train the denoising score function $\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t)$ for each noise level σ_t using a unified weighted objective:

$$\sum_{t=1}^T \sigma_t^2 \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x})} \|\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Noise-Conditioned Score Network (NCSN)

Train the denoising score function $\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t)$ for each noise level σ_t using a unified weighted objective:

$$\sum_{t=1}^T \sigma_t^2 \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x})} \|\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Here, $\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2} = -\frac{\epsilon}{\sigma_t}$

Noise-Conditioned Score Network (NCSN)

Train the denoising score function $\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t)$ for each noise level σ_t using a unified weighted objective:

$$\sum_{t=1}^T \sigma_t^2 \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x})} \|\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Here, $\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2} = -\frac{\epsilon}{\sigma_t}$

Training

1. Sample $\mathbf{x}_0 \sim p_{\text{data}}(\mathbf{x})$, $t \sim U\{1, T\}$, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$.
2. Compute noisy image $\mathbf{x}_t = \mathbf{x}_0 + \sigma_t \epsilon$.
3. Compute loss $\mathcal{L} = \sigma_t^2 \left\| \mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t) + \frac{\epsilon}{\sigma_t} \right\|^2$.
4. Update θ via stochastic gradient descent.

Noise-Conditioned Score Network (NCSN)

Train the denoising score function $\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t)$ for each noise level σ_t using a unified weighted objective:

$$\sum_{t=1}^T \sigma_t^2 \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x})} \|\mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x})\|_2^2 \rightarrow \min_{\theta}$$

Here, $\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2} = -\frac{\epsilon}{\sigma_t}$

Training

1. Sample $\mathbf{x}_0 \sim p_{\text{data}}(\mathbf{x})$, $t \sim U\{1, T\}$, $\epsilon \sim \mathcal{N}(0, \mathbf{I})$.
2. Compute noisy image $\mathbf{x}_t = \mathbf{x}_0 + \sigma_t \epsilon$.
3. Compute loss $\mathcal{L} = \sigma_t^2 \left\| \mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t) + \frac{\epsilon}{\sigma_t} \right\|^2$.
4. Update θ via stochastic gradient descent.

How do we sample from such a model?

Noise-Conditioned Score Network (NCSN)

Sampling (Annealed Langevin Dynamics)

1. Sample $\mathbf{x}_T^0 \sim \mathcal{N}(0, \sigma_T^2 \mathbf{I}) \approx q(\mathbf{x}_T)$.
2. Update $\mathbf{x}_t^l$ via L steps of Langevin dynamics at each noise level σ_t :

$$\mathbf{x}_t^l = \mathbf{x}_t^{l-1} + \frac{\eta_t}{2} \mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t^{l-1}) + \sqrt{\eta_t} \boldsymbol{\epsilon}_t^l.$$

3. Update $\mathbf{x}_{t-1}^0 := \mathbf{x}_t^L$ and proceed to the next noise level σ_{t-1} .

Noise-Conditioned Score Network (NCSN)

Sampling (Annealed Langevin Dynamics)

1. Sample $\mathbf{x}_T^0 \sim \mathcal{N}(0, \sigma_T^2 \mathbf{I}) \approx q(\mathbf{x}_T)$.
2. Update $\mathbf{x}_t^l$ via L steps of Langevin dynamics at each noise level σ_t :

$$\mathbf{x}_t^l = \mathbf{x}_t^{l-1} + \frac{\eta_t}{2} \mathbf{s}_{\theta, \sigma_t}(\mathbf{x}_t^{l-1}) + \sqrt{\eta_t} \boldsymbol{\epsilon}_t^l.$$

3. Update $\mathbf{x}_{t-1}^0 := \mathbf{x}_t^L$ and proceed to the next noise level σ_{t-1} .



Outline


01 Score Matching

Denoising Score Matching

Noise-Conditioned Score Network (NCSN)

02 **Forward Gaussian Diffusion Process**

Generative Models Taxonomy



Forward Gaussian Diffusion Process

Let $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$, $\beta_t \ll 1$. Define a Markov chain:

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$

Forward Gaussian Diffusion Process

Let $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$, $\beta_t \ll 1$. Define a Markov chain:

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I}) \\ q(\mathbf{x}_t | \mathbf{x}_{t-1}) &= \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})\end{aligned}$$

Forward Gaussian Diffusion Process

Let $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$, $\beta_t \ll 1$. Define a Markov chain:

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I}) \\ q(\mathbf{x}_t | \mathbf{x}_{t-1}) &= \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})\end{aligned}$$

Langevin Dynamics

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_{\theta}(\mathbf{x}_l) + \sqrt{\eta} \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(0, \mathbf{I})$$

Forward Gaussian Diffusion Process

Let $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$, $\beta_t \ll 1$. Define a Markov chain:

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I}) \\ q(\mathbf{x}_t | \mathbf{x}_{t-1}) &= \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})\end{aligned}$$

Langevin Dynamics

$$\begin{aligned}\mathbf{x}_{l+1} &= \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_{\theta}(\mathbf{x}_l) + \sqrt{\eta} \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(0, \mathbf{I}) \\ \mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t \\ &\approx \left(1 - \frac{\beta_t}{2}\right) \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t \\ &= \mathbf{x}_{t-1} + \frac{\beta_t}{2} (-\mathbf{x}_{t-1}) + \sqrt{\beta_t} \boldsymbol{\epsilon}_t\end{aligned}$$

Forward Gaussian Diffusion Process

Let $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$, $\beta_t \ll 1$. Define a Markov chain:

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I}) \\ q(\mathbf{x}_t | \mathbf{x}_{t-1}) &= \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})\end{aligned}$$

Langevin Dynamics

$$\begin{aligned}\mathbf{x}_{l+1} &= \mathbf{x}_l + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}_l} \log p_{\theta}(\mathbf{x}_l) + \sqrt{\eta} \boldsymbol{\epsilon}_l, \quad \boldsymbol{\epsilon}_l \sim \mathcal{N}(0, \mathbf{I}) \\ \mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t \\ &\approx \left(1 - \frac{\beta_t}{2}\right) \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t \\ &= \mathbf{x}_{t-1} + \frac{\beta_t}{2} (-\mathbf{x}_{t-1}) + \sqrt{\beta_t} \boldsymbol{\epsilon}_t\end{aligned}$$

- $\beta_t = \eta$
- $\nabla_{\mathbf{x}_{t-1}} \log p_{\theta}(\mathbf{x}_{t-1}) = -\mathbf{x}_{t-1} = \nabla_{\mathbf{x}_{t-1}} \log \mathcal{N}(0, \mathbf{I})$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$
$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$
$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Statement 1

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s = \prod_{s=1}^t (1 - \beta_s)$. Then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$
$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Statement 1

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s = \prod_{s=1}^t (1 - \beta_s)$. Then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Thus, samples at any timestep t can be generated directly from $\mathbf{x}_0$

$$\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t$$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Statement 1

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s = \prod_{s=1}^t (1 - \beta_s)$. Then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Thus, samples at any timestep t can be generated directly from $\mathbf{x}_0$

$$\begin{aligned} \mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-1}) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \end{aligned}$$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Statement 1

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s = \prod_{s=1}^t (1 - \beta_s)$. Then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Thus, samples at any timestep t can be generated directly from $\mathbf{x}_0$

$$\begin{aligned} \mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-1}) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + (\sqrt{\alpha_t (1 - \alpha_{t-1})} \boldsymbol{\epsilon}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t) \end{aligned}$$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Statement 1

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s = \prod_{s=1}^t (1 - \beta_s)$. Then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Thus, samples at any timestep t can be generated directly from $\mathbf{x}_0$

$$\begin{aligned} \mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-1}) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + (\sqrt{\alpha_t (1 - \alpha_{t-1})} \boldsymbol{\epsilon}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t) \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \boldsymbol{\epsilon}'_t \end{aligned}$$

Forward Gaussian Diffusion Process

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \boldsymbol{\epsilon}_t, \quad \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$$

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

Statement 1

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s = \prod_{s=1}^t (1 - \beta_s)$. Then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

Thus, samples at any timestep t can be generated directly from $\mathbf{x}_0$


$$\begin{aligned} \mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-1}) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + (\sqrt{\alpha_t (1 - \alpha_{t-1})} \boldsymbol{\epsilon}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t) \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \boldsymbol{\epsilon}'_t \\ &= \dots = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}) \end{aligned}$$

Forward Gaussian Diffusion Process

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N} \left(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I} \right);$$

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N} \left(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I} \right)$$

$$p_{\text{data}}(x_0) = \frac{1}{2} \mathcal{N}(-2, 0.35^2) + \frac{1}{2} \mathcal{N}(2, 0.35^2), \quad \beta_t = 0.01$$



The plots show marginal densities; the kernels above are conditional.

Forward Gaussian Diffusion Process

Statement 2

Applying the Markov chain to any distribution $p_{\text{data}}(\mathbf{x})$ yields $\mathbf{x}_{\infty} \sim p_{\infty}(\mathbf{x}) = \mathcal{N}(0, \mathbf{I})$, the **stationary** (limiting) distribution:

$$p_{\infty}(\mathbf{x}) = \int q(\mathbf{x}|\mathbf{x}')p_{\infty}(\mathbf{x}')d\mathbf{x}'$$

Forward Gaussian Diffusion Process

Statement 2

Applying the Markov chain to any distribution $p_{\text{data}}(\mathbf{x})$ yields $\mathbf{x}_{\infty} \sim p_{\infty}(\mathbf{x}) = \mathcal{N}(0, \mathbf{I})$, the **stationary** (limiting) distribution:

$$\begin{aligned} p_{\infty}(\mathbf{x}) &= \int q(\mathbf{x}|\mathbf{x}')p_{\infty}(\mathbf{x}')d\mathbf{x}' \\ p_{\infty}(\mathbf{x}) &= \int q(\mathbf{x}_{\infty}|\mathbf{x}_0)p_{\text{data}}(\mathbf{x}_0)d\mathbf{x}_0 \\ &\approx \mathcal{N}(0, \mathbf{I}) \int p_{\text{data}}(\mathbf{x}_0)d\mathbf{x}_0 = \mathcal{N}(0, \mathbf{I}) \end{aligned}$$

Forward Gaussian Diffusion Process

Statement 2


Applying the Markov chain to any distribution $p_{\text{data}}(\mathbf{x})$ yields $\mathbf{x}_\infty \sim p_\infty(\mathbf{x}) = \mathcal{N}(0, \mathbf{I})$, the **stationary** (limiting) distribution:

$$\begin{aligned} p_\infty(\mathbf{x}) &= \int q(\mathbf{x}|\mathbf{x}')p_\infty(\mathbf{x}')d\mathbf{x}' \\ p_\infty(\mathbf{x}) &= \int q(\mathbf{x}_\infty|\mathbf{x}_0)p_{\text{data}}(\mathbf{x}_0)d\mathbf{x}_0 \\ &\approx \mathcal{N}(0, \mathbf{I}) \int p_{\text{data}}(\mathbf{x}_0)d\mathbf{x}_0 = \mathcal{N}(0, \mathbf{I}) \end{aligned}$$

Note: This holds iff $\bar{\alpha}_t \rightarrow 0$, i.e., $\sum_{t=1}^{\infty} \beta_t = +\infty$.


Forward Gaussian Diffusion Process

Diffusion describes the migration of particles from regions of high density to those of low density.



Forward Gaussian Diffusion Process


Diffusion describes the migration of particles from regions of high density to those of low density.



1. $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$
2. $\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \epsilon_t,$
 $\epsilon_t \sim \mathcal{N}(0, \mathbf{I}), t \geq 1$
3. After $T \gg 1$ steps: $\mathbf{x}_T \sim p_{\infty}(\mathbf{x}) = \mathcal{N}(0, \mathbf{I})$

Forward Gaussian Diffusion Process

Diffusion describes the migration of particles from regions of high density to those of low density.



1. $\mathbf{x}_0 = \mathbf{x} \sim p_{\text{data}}(\mathbf{x})$
2. $\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \epsilon_t,$
 $\epsilon_t \sim \mathcal{N}(0, \mathbf{I}), t \geq 1$
3. After $T \gg 1$ steps: $\mathbf{x}_T \sim p_{\infty}(\mathbf{x}) = \mathcal{N}(0, \mathbf{I})$

If this process can be reversed, we can sample from $p_{\text{data}}(\mathbf{x})$ by starting from noise $p_{\infty}(\mathbf{x}) = \mathcal{N}(0, \mathbf{I})$. Our goal now becomes inverting this diffusion.

Denoising Score Matching

NCSN

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_0, \sigma_t^2 \mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \sigma_T^2 \mathbf{I})$$
$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2}$$

Denoising Score Matching

NCSN

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_0, \sigma_t^2 \mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \sigma_T^2 \mathbf{I})$$
$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2}$$

Gaussian Diffusion

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \mathbf{I})$$
$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0) = -\frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0}{1 - \bar{\alpha}_t}$$

Denoising Score Matching

NCSN

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_0, \sigma_t^2 \mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \sigma_T^2 \mathbf{I})$$

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2}$$

Gaussian Diffusion

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \mathbf{I})$$

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0) = -\frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t}\mathbf{x}_0}{1 - \bar{\alpha}_t}$$

Theorem (Denoising Score Matching)

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_t)} \|\mathbf{s}_{\theta,t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t)\|_2^2 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x})} \|\mathbf{s}_{\theta,t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x})\|_2^2 + \text{const}(\boldsymbol{\theta}) \end{aligned}$$

Denoising Score Matching

NCSN

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_0, \sigma_t^2 \mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \sigma_T^2 \mathbf{I})$$

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}) = -\frac{\mathbf{x}_t - \mathbf{x}}{\sigma_t^2}$$

Gaussian Diffusion

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}), \quad q(\mathbf{x}_1) \approx p_{\text{data}}(\mathbf{x}), \quad q(\mathbf{x}_T) \approx \mathcal{N}(0, \mathbf{I})$$

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0) = -\frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0}{1 - \bar{\alpha}_t}$$

Theorem (Denoising Score Matching)

$$\begin{aligned} & \mathbb{E}_{q(\mathbf{x}_t)} \|\mathbf{s}_{\theta,t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t)\|_2^2 \\ &= \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x})} \|\mathbf{s}_{\theta,t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x})\|_2^2 + \text{const}(\boldsymbol{\theta}) \end{aligned}$$

Note: Annealed Langevin dynamics applies to diffusion, too.

NCSN with annealed Langevin dynamics is an early score-based generative framework (2019–2020).

Summary

- Score matching minimizes Fisher divergence to estimate the score function.
- Denoising score matching makes the score objective estimable using corrupted samples and the known corruption kernel.
- Tweedie's formula converts the score into MSE-optimal clean-data and noise predictions.
- Noise-Conditioned Score Networks learn scores at multiple noise levels and sample with annealed Langevin dynamics.
- The forward Gaussian diffusion process admits a closed-form conditional distribution for any timestep.
- With a suitable noise schedule, the forward process converges to standard Gaussian noise; denoising score matching also applies to these diffusion marginals.