

MIPT & YSDA · AUTUMN 2026

Deep Generative Models

Lecture 10

Roman Isachenko

Moscow Institute of Physics and Technology
Yandex School of Data Analysis

Recap of Previous Lecture

Continuous-Time Dynamics

$$\frac{d\mathbf{x}(t)}{dt} = \mathbf{v}_{\boldsymbol{\theta}}(\mathbf{x}(t), t); \quad \text{with initial condition } \mathbf{x}(t_0) = \mathbf{x}_0.$$
$$\boldsymbol{\psi}_t(\mathbf{x}_0) = \int_{t_0}^t \mathbf{v}_{\boldsymbol{\theta}}(\mathbf{x}(s), s) ds + \mathbf{x}_0 \approx \text{ODESolve}_v(\mathbf{x}_0, \boldsymbol{\theta}, t_0, t).$$

- $\mathbf{v}_{\boldsymbol{\theta}} : \mathbb{R}^m \times [t_0, t_1] \rightarrow \mathbb{R}^m$ is a vector field.
- $\boldsymbol{\psi} : \mathbb{R}^m \times [t_0, t_1] \rightarrow \mathbb{R}^m$ is a flow (the solution of ODE):

Euler ODESolve


$$\mathbf{x}(t + h) = \mathbf{x}(t) + h \cdot \mathbf{v}_{\boldsymbol{\theta}}(\mathbf{x}(t), t)$$

More advanced numerical methods (such as Runge-Kutta) are often used instead of unstable Euler update step.

Recap of Previous Lecture

$$\frac{d\mathbf{x}(t)}{dt} = \mathbf{v}_\theta(\mathbf{x}(t), t); \quad \mathbf{x}(t_0) = \mathbf{x}_0$$

- Suppose $\mathbf{x}(0)$ is a random variable with density $p_0(\mathbf{x})$. Then, $\mathbf{x}(t)$ is a random variable with density $p_t(\mathbf{x})$.
- $p_t(\mathbf{x}) = p(\mathbf{x}, t)$ describes the **probability path** between $p_0(\mathbf{x})$ and $p_1(\mathbf{x})$.



Recap of Previous Lecture

Theorem (Picard)

If $\mathbf{v}$ is continuously differentiable with a bounded derivative in $\mathbf{x}$ and continuous in t , then the ODE has a **unique solution** given by a flow ψ_t .

This guarantees the ODE is **uniquely reversible**.

$$\psi_1(\mathbf{x}_0) = \mathbf{x}_0 + \int_0^1 \mathbf{v}_\theta(\psi_t(\mathbf{x}_0), t) dt$$

$$\mathbf{x}(1) = \mathbf{x}(0) + \int_0^1 \mathbf{v}_\theta(\mathbf{x}(t), t) dt$$

$$\mathbf{x}(0) = \mathbf{x}(1) + \int_1^0 \mathbf{v}_\theta(\mathbf{x}(t), t) dt$$

Note: Unlike discrete-time flows, $\mathbf{v}$ need not be invertible (uniqueness ensures bijection).

How can we compute $p_t(\mathbf{x})$ at arbitrary t ?

Recap of Previous Lecture

$$\frac{d\mathbf{x}(t)}{dt} = \mathbf{v}_\theta(\mathbf{x}(t), t); \quad \text{with initial condition } \mathbf{x}(t_0) = \mathbf{x}_0$$

Theorem (Continuity Equation)

If $\mathbf{v}$ is continuously differentiable in $\mathbf{x}$ (with bounded derivative) and continuous in t , and $p_t(\mathbf{x}) > 0$, then

$$\frac{d \log p_t(\mathbf{x}(t))}{dt} = - \operatorname{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right)$$

Solution

$$\log p_1(\mathbf{x}(1)) = \log p_0(\mathbf{x}(0)) - \int_0^1 \operatorname{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right) dt.$$

- This solution gives us the density **along the trajectory**.
- However, it's difficult to efficiently estimate **the last term**.

Recap of Previous Lecture

SDE Basics

Let's define a stochastic process $\mathbf{x}(t)$ with initial condition $\mathbf{x}(0) \sim p_0(\mathbf{x})$:

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w},$$

where $\mathbf{w}(t)$ is the standard Wiener process (Brownian motion):

$$\mathbf{w}(t) - \mathbf{w}(s) \sim \mathcal{N}(0, (t - s)\mathbf{I}), \quad d\mathbf{w} = \boldsymbol{\epsilon} \cdot \sqrt{dt}, \text{ where } \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I}).$$

Discretizing the SDE (Euler-Maruyama Update) – SDEsolve

$$\mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{f}(\mathbf{x}(t), t) \cdot dt + g(t) \cdot \boldsymbol{\epsilon} \cdot \sqrt{dt}$$

- At each time t , we have the density $p_t(\mathbf{x}) = p(\mathbf{x}, t)$.
- $p : \mathbb{R}^m \times [0, 1] \rightarrow \mathbb{R}_+$ is a **probability path** between $p_0(\mathbf{x})$ and $p_1(\mathbf{x})$.

Recap of Previous Lecture

Theorem (Kolmogorov-Fokker-Planck)

If $p_t(\mathbf{x}) \in C^{1,2}$ (i.e., C^1 in t and C^2 in $\mathbf{x}$), then

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})) + \frac{1}{2}g^2(t)\Delta_{\mathbf{x}}p_t(\mathbf{x})$$

Langevin SDE (Special Case)

$$d\mathbf{x} = \frac{1}{2} \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) dt + \mathbf{1} \cdot d\mathbf{w}$$

The density $p_t(\mathbf{x})$ is a **stationary** distribution for the SDE. Setting the stationary density to the model distribution, $p_t(\mathbf{x}) = p_{\theta}(\mathbf{x})$ for all t :

Langevin Dynamic

$$\mathbf{x}_{t+1} = \mathbf{x}_t + \frac{\eta}{2} \cdot \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) + \sqrt{\eta} \cdot \epsilon, \quad \eta \approx dt.$$

Outline

- 01 Diffusion and Score Matching SDEs (VP-SDE / VE-SDE)
- 02 Probability Flow ODE
- 03 Reverse SDE
- 04 Score-Based Generative Models Through SDEs
- 05 Flow Matching (FM)

Outline

01 Diffusion and Score Matching SDEs (VP-SDE / VE-SDE)

02 Probability Flow ODE

03 Reverse SDE

04 Score-Based Generative Models Through SDEs

05 Flow Matching (FM)

Score Matching SDE

Denoising Score Matching

$$\begin{aligned}\mathbf{x}_t &= \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, & q(\mathbf{x}_t | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I}) \\ \mathbf{x}_{t-1} &= \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, & q(\mathbf{x}_{t-1} | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})\end{aligned}$$

Score Matching SDE

Denoising Score Matching

$$\begin{aligned}\mathbf{x}_t &= \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, & q(\mathbf{x}_t | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I}) \\ \mathbf{x}_{t-1} &= \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, & q(\mathbf{x}_{t-1} | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I}) \\ \mathbf{x}_t &= \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, & q(\mathbf{x}_t | \mathbf{x}_{t-1}) &= \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})\end{aligned}$$

Score Matching SDE

Denoising Score Matching

$$\begin{aligned}\mathbf{x}_t &= \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, & q(\mathbf{x}_t | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I}) \\ \mathbf{x}_{t-1} &= \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, & q(\mathbf{x}_{t-1} | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})\end{aligned}$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Let's transform this Markov chain into the continuous stochastic process $\mathbf{x}(t)$ by letting $T \rightarrow \infty$:

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\sigma^2(t) - \sigma^2(t - dt)} \cdot \boldsymbol{\epsilon}$$

Score Matching SDE

Denoising Score Matching

$$\begin{aligned}\mathbf{x}_t &= \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, & q(\mathbf{x}_t | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I}) \\ \mathbf{x}_{t-1} &= \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, & q(\mathbf{x}_{t-1} | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})\end{aligned}$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Let's transform this Markov chain into the continuous stochastic process $\mathbf{x}(t)$ by letting $T \rightarrow \infty$:

$$\begin{aligned}\mathbf{x}(t) &= \mathbf{x}(t - dt) + \sqrt{\sigma^2(t) - \sigma^2(t - dt)} \cdot \boldsymbol{\epsilon} \\ &= \mathbf{x}(t - dt) + \sqrt{\frac{\sigma^2(t) - \sigma^2(t - dt)}{dt}} dt \cdot \boldsymbol{\epsilon}\end{aligned}$$

Score Matching SDE

Denoising Score Matching

$$\begin{aligned}\mathbf{x}_t &= \mathbf{x} + \sigma_t \cdot \boldsymbol{\epsilon}_t, & q(\mathbf{x}_t | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_t^2 \cdot \mathbf{I}) \\ \mathbf{x}_{t-1} &= \mathbf{x} + \sigma_{t-1} \cdot \boldsymbol{\epsilon}_{t-1}, & q(\mathbf{x}_{t-1} | \mathbf{x}) &= \mathcal{N}(\mathbf{x}, \sigma_{t-1}^2 \cdot \mathbf{I})\end{aligned}$$

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2} \cdot \boldsymbol{\epsilon}, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2) \cdot \mathbf{I})$$

Let's transform this Markov chain into the continuous stochastic process $\mathbf{x}(t)$ by letting $T \rightarrow \infty$:

$$\begin{aligned}\mathbf{x}(t) &= \mathbf{x}(t - dt) + \sqrt{\sigma^2(t) - \sigma^2(t - dt)} \cdot \boldsymbol{\epsilon} \\ &= \mathbf{x}(t - dt) + \sqrt{\frac{\sigma^2(t) - \sigma^2(t - dt)}{dt}} dt \cdot \boldsymbol{\epsilon} \\ &= \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}\end{aligned}$$

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$
$$d\mathbf{x} = \mathbf{x}(t) - \mathbf{x}(t - dt)$$

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$
$$d\mathbf{x} = \mathbf{x}(t) - \mathbf{x}(t - dt)$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$$d\mathbf{x} = \mathbf{x}(t) - \mathbf{x}(t - dt)$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$$\mathbf{f}(\mathbf{x}, t) = 0, \quad g(t) = \sqrt{\frac{d[\sigma^2(t)]}{dt}}$$

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}_t, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}_t, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \rightarrow \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\mathbf{x}(t) = \sqrt{1 - \beta(t)dt} \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon}$$

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}_t, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \rightarrow \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\begin{aligned} \mathbf{x}(t) &= \sqrt{1 - \beta(t)dt} \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} \\ &\approx \left(1 - \frac{1}{2}\beta(t)dt\right) \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} \end{aligned}$$

Diffusion SDE

Denoising Diffusion

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}_t, \quad q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1}, \beta_t \cdot \mathbf{I})$$

Let's turn this Markov chain into a continuous stochastic process by letting $T \rightarrow \infty$ and setting $\beta_t = \beta(\frac{t}{T}) \cdot \frac{1}{T}$ (where $dt = \frac{1}{T}$):

$$\begin{aligned} \mathbf{x}(t) &= \sqrt{1 - \beta(t)dt} \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)dt} \cdot \boldsymbol{\epsilon} \\ &\approx \left(1 - \frac{1}{2}\beta(t)dt\right) \cdot \mathbf{x}(t - dt) + \sqrt{\beta(t)}d\mathbf{w} \cdot \boldsymbol{\epsilon} \\ &= \mathbf{x}(t - dt) - \frac{1}{2}\beta(t)\mathbf{x}(t - dt)dt + \sqrt{\beta(t)} \cdot d\mathbf{w} \end{aligned}$$

Diffusion SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) - \frac{1}{2}\beta(t)\mathbf{x}(t - dt)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

Diffusion SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) - \frac{1}{2}\beta(t)\mathbf{x}(t - dt)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

Variance Preserving SDE

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$
$$\mathbf{f}(\mathbf{x}, t) = -\frac{1}{2}\beta(t)\mathbf{x}(t), \quad g(t) = \sqrt{\beta(t)}$$

Variance is preserved as long as $\mathbf{x}(0)$ has unit variance.

Diffusion SDE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

Diffusion SDE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

Variance Exploding SDE (NCSN)

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Preserving SDE (DDPM)

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

We treat the discrete-time methods (NCSN, DDPM) as a special case of the continuous-time methods (VE-SDE, VP-SDE).

Outline

01 Diffusion and Score Matching SDEs (VP-SDE / VE-SDE)

02 Probability Flow ODE

03 Reverse SDE

04 Score-Based Generative Models Through SDEs

05 Flow Matching (FM)

Probability Flow ODE

ODE and Continuity Equation

$$\begin{aligned} d\mathbf{x} &= \mathbf{v}(\mathbf{x}, t)dt \\ \frac{d \log p_t(\mathbf{x}(t))}{dt} &= -\operatorname{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right) \Leftrightarrow \frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x})) \end{aligned}$$

The only source of randomness is the initial distribution $p_0(\mathbf{x})$.

Probability Flow ODE

ODE and Continuity Equation

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt$$

$$\frac{d \log p_t(\mathbf{x}(t))}{dt} = -\operatorname{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right) \Leftrightarrow \frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x}))$$

The only source of randomness is the initial distribution $p_0(\mathbf{x})$.

SDE and KFP Equation

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})) + \frac{1}{2}g^2(t)\Delta_{\mathbf{x}}p_t(\mathbf{x})$$

Now there are two sources of randomness: the initial distribution $p_0(\mathbf{x})$ and the Wiener process $\mathbf{w}(t)$.

Probability Flow ODE

Theorem

Suppose the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ induces the probability path $p_t(\mathbf{x})$ with $p_t(\mathbf{x}) > 0$. Then, there exists an ODE with the same probability path $p_t(\mathbf{x})$, given by


$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Probability Flow ODE

Theorem

Suppose the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ induces the probability path $p_t(\mathbf{x})$ with $p_t(\mathbf{x}) > 0$. Then, there exists an ODE with the same probability path $p_t(\mathbf{x})$, given by

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$



Probability Flow ODE

Theorem

Suppose the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ induces the probability path $p_t(\mathbf{x})$ with $p_t(\mathbf{x}) > 0$. Then, there exists an ODE with the same probability path $p_t(\mathbf{x})$, given by

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Proof

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})] + \frac{1}{2}g^2(t) \frac{\partial^2 p_t(\mathbf{x})}{\partial \mathbf{x}^2} \right)$$

Probability Flow ODE

Theorem

Suppose the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ induces the probability path $p_t(\mathbf{x})$ with $p_t(\mathbf{x}) > 0$. Then, there exists an ODE with the same probability path $p_t(\mathbf{x})$, given by

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Proof

$$\begin{aligned} \frac{\partial p_t(\mathbf{x})}{\partial t} &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})] + \frac{1}{2}g^2(t) \frac{\partial^2 p_t(\mathbf{x})}{\partial \mathbf{x}^2} \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x}) - \frac{1}{2}g^2(t) \frac{\partial p_t(\mathbf{x})}{\partial \mathbf{x}} \right] \right) \end{aligned}$$

Probability Flow ODE

Theorem

Suppose the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ induces the probability path $p_t(\mathbf{x})$ with $p_t(\mathbf{x}) > 0$. Then, there exists an ODE with the same probability path $p_t(\mathbf{x})$, given by

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Proof

$$\begin{aligned} \frac{\partial p_t(\mathbf{x})}{\partial t} &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})] + \frac{1}{2}g^2(t) \frac{\partial^2 p_t(\mathbf{x})}{\partial \mathbf{x}^2} \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x}) - \frac{1}{2}g^2(t) \frac{\partial p_t(\mathbf{x})}{\partial \mathbf{x}} \right] \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x}) - \frac{1}{2}g^2(t)p_t(\mathbf{x}) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right] \right) \end{aligned}$$

Probability Flow ODE

Theorem

Suppose the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ induces the probability path $p_t(\mathbf{x})$ with $p_t(\mathbf{x}) > 0$. Then, there exists an ODE with the same probability path $p_t(\mathbf{x})$, given by

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Proof

$$\begin{aligned} \frac{\partial p_t(\mathbf{x})}{\partial t} &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x})] + \frac{1}{2}g^2(t) \frac{\partial^2 p_t(\mathbf{x})}{\partial \mathbf{x}^2} \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x}) - \frac{1}{2}g^2(t) \frac{\partial p_t(\mathbf{x})}{\partial \mathbf{x}} \right] \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\mathbf{f}(\mathbf{x}, t)p_t(\mathbf{x}) - \frac{1}{2}g^2(t) p_t(\mathbf{x}) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right] \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) p_t(\mathbf{x}) \right] \right) \end{aligned}$$

Probability Flow ODE

Proof (Continued)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) p_t(\mathbf{x}) \right] \right)$$

Probability Flow ODE

Proof (Continued)

$$\begin{aligned}\frac{\partial p_t(\mathbf{x})}{\partial t} &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) p_t(\mathbf{x}) \right] \right) \\ &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} [\mathbf{v}(\mathbf{x}, t) p_t(\mathbf{x})] \right) = -\text{div} (\mathbf{v}(\mathbf{x}, t) p_t(\mathbf{x}))\end{aligned}$$

Probability Flow ODE

Proof (Continued)

$$\begin{aligned}
 \frac{\partial p_t(\mathbf{x})}{\partial t} &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} \left[\left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) p_t(\mathbf{x}) \right] \right) \\
 &= \text{tr} \left(-\frac{\partial}{\partial \mathbf{x}} [\mathbf{v}(\mathbf{x}, t) p_t(\mathbf{x})] \right) = -\text{div}(\mathbf{v}(\mathbf{x}, t) p_t(\mathbf{x})) \\
 \mathbf{v}(\mathbf{x}, t) &= \mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}); \quad \tilde{g}(t) = 0 \\
 d\mathbf{x} &= \mathbf{v}(\mathbf{x}, t) dt + 0 \cdot d\mathbf{w} = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt \\
 \frac{\partial p_t(\mathbf{x})}{\partial t} &= -\text{div}(\mathbf{v}(\mathbf{x}, t) p_t(\mathbf{x})) + \frac{1}{2} \tilde{g}^2(t) \Delta_{\mathbf{x}} p_t(\mathbf{x})
 \end{aligned}$$

Probability Flow ODE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w} \quad - \text{SDE}$$

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt \quad - \text{Probability Flow ODE}$$

$$\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$$

Probability Flow ODE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w} \quad - \text{SDE}$$

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt \quad - \text{Probability Flow ODE}$$

$$\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$$

- $\mathbf{s}(\mathbf{x}, t) = \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$ is the continuous-time score function.


Probability Flow ODE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w} \quad - \text{SDE}$$

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt \quad - \text{Probability Flow ODE}$$

$$\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$$

- $\mathbf{s}(\mathbf{x}, t) = \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$ is the continuous-time score function.
- The ODE produces more stable trajectories.



Why convert SDEs to PF-ODEs? (1) Efficient Sampling

- Brownian increments scale as $d\mathbf{w} \sim \sqrt{dt}$ — SDE trajectories are nowhere differentiable.
- `SDESo1ve` (Euler–Maruyama): local error $\mathcal{O}(\sqrt{\Delta t})$.
- Higher-order SDE schemes need Lévy areas $\Rightarrow$ impractical in high dimensions.

Why convert SDEs to PF-ODEs? (1) Efficient Sampling

- Brownian increments scale as $d\mathbf{w} \sim \sqrt{dt}$ — SDE trajectories are nowhere differentiable.
- `SDESo1ve` (Euler–Maruyama): local error $\mathcal{O}(\sqrt{\Delta t})$.
- Higher-order SDE schemes need Lévy areas $\Rightarrow$ impractical in high dimensions.

PF-ODE drift $\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t)$ **is smooth**

- Euler: $\mathcal{O}(\Delta t)$, Heun: $\mathcal{O}(\Delta t^2)$, RK4: $\mathcal{O}(\Delta t^4)$.
- Adaptive step-size control is standard for `ODESo1ve`.

Why convert SDEs to PF-ODEs? (1) Efficient Sampling

- Brownian increments scale as $d\mathbf{w} \sim \sqrt{dt}$ — SDE trajectories are nowhere differentiable.
- `SDESo1ve` (Euler–Maruyama): local error $\mathcal{O}(\sqrt{\Delta t})$.
- Higher-order SDE schemes need Lévy areas $\Rightarrow$ impractical in high dimensions.

PF-ODE drift $\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t)$ is smooth

- Euler: $\mathcal{O}(\Delta t)$, Heun: $\mathcal{O}(\Delta t^2)$, RK4: $\mathcal{O}(\Delta t^4)$.
- Adaptive step-size control is standard for `ODESo1ve`.
- In practice: 100–1000 `SDESo1ve` steps $\rightarrow$ 20–50 `ODESo1ve` steps.
- Caveat: SDE sampling can still win at high NFE — injected noise corrects accumulated error.

Why convert SDEs to PF-ODEs? (2) Deterministic Sampling

- The PF-ODE defines a deterministic, invertible map $\mathbf{x}(0) \leftrightarrow \mathbf{x}(T)$.
- Each terminal $\mathbf{x}(T)$ produces a unique trajectory back to data space.

Why convert SDEs to PF-ODEs? (2) Deterministic Sampling

- The PF-ODE defines a deterministic, invertible map $\mathbf{x}(0) \leftrightarrow \mathbf{x}(T)$.
- Each terminal $\mathbf{x}(T)$ produces a unique trajectory back to data space.

Exact log-likelihood (instantaneous change of variables)

$$\log p_0(\mathbf{x}(0)) = \log p_T(\mathbf{x}(T)) + \int_0^T \text{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right) dt$$

Why convert SDEs to PF-ODEs? (2) Deterministic Sampling

- The PF-ODE defines a deterministic, invertible map $\mathbf{x}(0) \leftrightarrow \mathbf{x}(T)$.
- Each terminal $\mathbf{x}(T)$ produces a unique trajectory back to data space.

Exact log-likelihood (instantaneous change of variables)

$$\log p_0(\mathbf{x}(0)) = \log p_T(\mathbf{x}(T)) + \int_0^T \text{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right) dt$$

- Connects PF-ODE to continuous normalizing flows.
- Invertibility enables image $\rightarrow$ latent $\rightarrow$ edited image workflows.
- Smooth interpolation in latent space between two samples.

Sampling: Reverse SDE vs PF-ODE

Same trained score $s_\theta(\mathbf{x}, t)$; only the discretization changes (reverse SDE derived in the next section).

Sampling: Reverse SDE vs PF-ODE

Same trained score $s_\theta(\mathbf{x}, t)$; only the discretization changes (reverse SDE derived in the next section).

Sampling (Reverse SDE, Euler–Maruyama)

1. Sample $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$.
2. Initialize step $\Delta t = T/K$.
3. Denoise:

$$\mathbf{x}_{t-\Delta t} = \mathbf{x}_t - \left[\mathbf{f}(\mathbf{x}_t, t) - g^2(t) \mathbf{s}_\theta(\mathbf{x}_t, t) \right] \Delta t + g(t) \sqrt{\Delta t} \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$$

Sampling: Reverse SDE vs PF-ODE

Same trained score $s_\theta(\mathbf{x}, t)$; only the discretization changes (reverse SDE derived in the next section).

Sampling (Reverse SDE, Euler–Maruyama)

1. Sample $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$.
2. Initialize step $\Delta t = T/K$.
3. Denoise:

$$\mathbf{x}_{t-\Delta t} = \mathbf{x}_t - \left[\mathbf{f}(\mathbf{x}_t, t) - g^2(t) s_\theta(\mathbf{x}_t, t) \right] \Delta t + g(t) \sqrt{\Delta t} \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$$

Sampling (PF-ODE: Euler / Heun / RK4)

1. Sample $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$.
2. Initialize step $\Delta t = T/K$ (with $K \ll$ SDE case).
3. Denoise:

$$\mathbf{x}_{t-\Delta t} = \mathbf{x}_t - \left[\mathbf{f}(\mathbf{x}_t, t) - \frac{1}{2} g^2(t) s_\theta(\mathbf{x}_t, t) \right] \Delta t \quad (\text{no noise injection})$$

Outline

- 01 Diffusion and Score Matching SDEs (VP-SDE / VE-SDE)
- 02 Probability Flow ODE
- 03 Reverse SDE**
- 04 Score-Based Generative Models Through SDEs
- 05 Flow Matching (FM)

Reverse SDE

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt, \quad \mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{v}(\mathbf{x}, t)dt$$

Here dt can be > 0 or < 0 .

Reverse SDE

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt, \quad \mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{v}(\mathbf{x}, t)dt$$

Here dt can be > 0 or < 0 .

Reverse ODE

Let $\tau = 1 - t$ ($d\tau = -dt$).

$$d\mathbf{x} = -\mathbf{v}(\mathbf{x}, 1 - \tau)d\tau$$

Reverse SDE

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt, \quad \mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{v}(\mathbf{x}, t)dt$$

Here dt can be > 0 or < 0 .

Reverse ODE

Let $\tau = 1 - t$ ($d\tau = -dt$).

$$d\mathbf{x} = -\mathbf{v}(\mathbf{x}, 1 - \tau)d\tau$$

- How do we reverse the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$?
- The Wiener process introduces randomness that must be reversed.

Reverse SDE

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt, \quad \mathbf{x}(t + dt) = \mathbf{x}(t) + \mathbf{v}(\mathbf{x}, t)dt$$

Here dt can be > 0 or < 0 .

Reverse ODE

Let $\tau = 1 - t$ ($d\tau = -dt$).

$$d\mathbf{x} = -\mathbf{v}(\mathbf{x}, 1 - \tau)d\tau$$

- How do we reverse the SDE $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$?
- The Wiener process introduces randomness that must be reversed.

Theorem

There exists a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (assume $p_t(\mathbf{x}) > 0$), given by:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

where $dt < 0$ and $\bar{\mathbf{w}}$ is a reverse-time Wiener process.

Reverse SDE

Theorem

There exists a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (assume $p_t(\mathbf{x}) > 0$), given by:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

where $dt < 0$ and $\bar{\mathbf{w}}(t)$ is a reverse-time Wiener process.

Reverse SDE

Theorem

There exists a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (assume $p_t(\mathbf{x}) > 0$), given by:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

where $dt < 0$ and $\bar{\mathbf{w}}(t)$ is a reverse-time Wiener process.

- Here the score function appears: $\mathbf{s}(\mathbf{x}, t) = \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$.
- Since $\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t)$ (PF-ODE), the reverse drift equals $\mathbf{f} - g^2\mathbf{s} = \mathbf{v} - \frac{1}{2}g^2\mathbf{s}$.

Reverse SDE

Theorem

There exists a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (assume $p_t(\mathbf{x}) > 0$), given by:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

where $dt < 0$ and $\bar{\mathbf{w}}(t)$ is a reverse-time Wiener process.

- Here the score function appears: $\mathbf{s}(\mathbf{x}, t) = \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x})$.
- Since $\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t)$ (PF-ODE), the reverse drift equals $\mathbf{f} - g^2\mathbf{s} = \mathbf{v} - \frac{1}{2}g^2\mathbf{s}$.

Proof Sketch

- Convert the initial SDE to a probability flow ODE.
- Reverse the probability flow ODE.
- Convert the reversed probability flow ODE back to an SDE.

Proof of Reverse SDE

- Convert the initial SDE to a probability flow ODE:

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

Proof of Reverse SDE

- Convert the initial SDE to a probability flow ODE:

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

- Reverse the probability flow ODE:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt$$

$$d\mathbf{x} = \left(-\mathbf{f}(\mathbf{x}, 1 - \tau) + \frac{1}{2}g^2(1 - \tau)\frac{\partial}{\partial \mathbf{x}} \log p_{1-\tau}(\mathbf{x}) \right) d\tau$$

Proof of Reverse SDE

- Convert the reversed probability flow ODE back to an SDE:

$$d\mathbf{x} = \left(-\mathbf{f}(\mathbf{x}, 1 - \tau) + \frac{1}{2}g^2(1 - \tau) \frac{\partial}{\partial \mathbf{x}} \log p_{1-\tau}(\mathbf{x}) \right) d\tau$$

$$d\mathbf{x} = \left(-\mathbf{f}(\mathbf{x}, 1 - \tau) + g^2(1 - \tau) \frac{\partial}{\partial \mathbf{x}} \log p_{1-\tau}(\mathbf{x}) \right) d\tau + g(1 - \tau)d\mathbf{w}$$

Here $\mathbf{w}(\tau)$ is a Wiener process in forward τ -time.

Reverse SDE

Theorem

There exists a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (assume $p_t(\mathbf{x}) > 0$), given by:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

where $dt < 0$ and $\bar{\mathbf{w}}(t)$ is a reverse-time Wiener process.

Proof (Continued)

$$d\mathbf{x} = \left(-\mathbf{f}(\mathbf{x}, 1 - \tau) + g^2(1 - \tau) \frac{\partial}{\partial \mathbf{x}} \log p_{1-\tau}(\mathbf{x}) \right) d\tau + g(1 - \tau)d\mathbf{w}$$

Reverse SDE

Theorem

There exists a reverse SDE for $d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$ (assume $p_t(\mathbf{x}) > 0$), given by:

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

where $dt < 0$ and $\bar{\mathbf{w}}(t)$ is a reverse-time Wiener process.

Proof (Continued)

$$d\mathbf{x} = \left(-\mathbf{f}(\mathbf{x}, 1 - \tau) + g^2(1 - \tau) \frac{\partial}{\partial \mathbf{x}} \log p_{1-\tau}(\mathbf{x}) \right) d\tau + g(1 - \tau)d\mathbf{w}$$

$$d\mathbf{x} = \left(\mathbf{f}(\mathbf{x}, t) - g^2(t) \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + g(t)d\bar{\mathbf{w}}$$

Here $d\tau > 0$, $dt < 0$, and $\bar{\mathbf{w}}$ is reverse-time Wiener process.

Reverse SDE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$$\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t), \quad \text{where} \quad \mathbf{s}(\mathbf{x}, t) = \nabla_{\mathbf{x}} \log p_t(\mathbf{x})$$

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt \quad - \text{Probability Flow ODE}$$

$$d\mathbf{x} = \left(\mathbf{v}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t) \right)dt + g(t)d\bar{\mathbf{w}} \quad - \text{Reverse SDE}$$


Reverse SDE

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}$$

$$\mathbf{v}(\mathbf{x}, t) = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t), \quad \text{where} \quad \mathbf{s}(\mathbf{x}, t) = \nabla_{\mathbf{x}} \log p_t(\mathbf{x})$$

$$d\mathbf{x} = \mathbf{v}(\mathbf{x}, t)dt \quad - \text{Probability Flow ODE}$$

$$d\mathbf{x} = \left(\mathbf{v}(\mathbf{x}, t) - \frac{1}{2}g^2(t)\mathbf{s}(\mathbf{x}, t) \right)dt + g(t)d\bar{\mathbf{w}} \quad - \text{Reverse SDE}$$



Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}, \quad \mathbf{f}(\mathbf{x}, t) = 0, \quad g(t) = \sqrt{\frac{d[\sigma^2(t)]}{dt}}$$

Score Matching SDE

$$\mathbf{x}(t) = \mathbf{x}(t - dt) + \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

Variance Exploding SDE

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w}$$

$\sigma(t)$ is a monotonically increasing function.

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t)dt + g(t)d\mathbf{w}, \quad \mathbf{f}(\mathbf{x}, t) = 0, \quad g(t) = \sqrt{\frac{d[\sigma^2(t)]}{dt}}$$

$$d\mathbf{x} = \left(-\frac{1}{2} \frac{d[\sigma^2(t)]}{dt} \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt \quad (\text{probability flow ODE})$$

$$d\mathbf{x} = \left(-\frac{d[\sigma^2(t)]}{dt} \frac{\partial}{\partial \mathbf{x}} \log p_t(\mathbf{x}) \right) dt + \sqrt{\frac{d[\sigma^2(t)]}{dt}} d\bar{\mathbf{w}} \quad (\text{reverse SDE})$$

Diffusion SDE

Variance Preserving SDE

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$
$$\mathbf{f}(\mathbf{x}, t) = -\frac{1}{2}\beta(t)\mathbf{x}(t), \quad g(t) = \sqrt{\beta(t)}$$

Variance is preserved as long as $\mathbf{x}(0)$ has unit variance.

Diffusion SDE

Variance Preserving SDE

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w}$$

$$\mathbf{f}(\mathbf{x}, t) = -\frac{1}{2}\beta(t)\mathbf{x}(t), \quad g(t) = \sqrt{\beta(t)}$$

Variance is preserved as long as $\mathbf{x}(0)$ has unit variance.

$$d\mathbf{x} = \left(-\frac{1}{2}\beta(t)\mathbf{x}(t) - \frac{1}{2}\beta(t)\frac{\partial}{\partial\mathbf{x}}\log p_t(\mathbf{x}) \right) dt \quad (\text{probability flow ODE})$$

$$d\mathbf{x} = \left(-\frac{1}{2}\beta(t)\mathbf{x}(t) - \beta(t)\frac{\partial}{\partial\mathbf{x}}\log p_t(\mathbf{x}) \right) dt + \sqrt{\beta(t)}d\bar{\mathbf{w}} \quad (\text{reverse SDE})$$

Outline

- 01 Diffusion and Score Matching SDEs (VP-SDE / VE-SDE)
- 02 Probability Flow ODE
- 03 Reverse SDE
- 04 Score-Based Generative Models Through SDEs**
- 05 Flow Matching (FM)

Score-Based Generative Models Through SDEs

Discrete-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}_0)} \mathbb{E}_{t \sim U\{1, T\}} \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} \left\| \mathbf{s}_{\theta, t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t | \mathbf{x}_0) \right\|_2^2$$

Is it possible to train score-based diffusion models in continuous time?

Score-Based Generative Models Through SDEs

Discrete-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}_0)} \mathbb{E}_{t \sim U\{1, T\}} \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} \left\| \mathbf{s}_{\theta, t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t | \mathbf{x}_0) \right\|_2^2$$

Is it possible to train score-based diffusion models in continuous time?

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0, 1]} \mathbb{E}_{q(\mathbf{x}(t) | \mathbf{x}(0))} \left\| \mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t) | \mathbf{x}(0)) \right\|_2^2$$

Score-Based Generative Models Through SDEs


Discrete-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}_0)} \mathbb{E}_{t \sim U\{1, T\}} \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} \left\| \mathbf{s}_{\theta, t}(\mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t | \mathbf{x}_0) \right\|_2^2$$


Is it possible to train score-based diffusion models in continuous time?

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0, 1]} \mathbb{E}_{q(\mathbf{x}(t) | \mathbf{x}(0))} \left\| \mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t) | \mathbf{x}(0)) \right\|_2^2$$



Generative Models Taxonomy



Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))} \left\| \mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) \right\|_2^2$$

Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))} \left\| \mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) \right\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

$$\nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) = -\frac{1}{\sigma^2(t)}(\mathbf{x}(t) - \boldsymbol{\mu})$$

Note: Normality holds for $\mathbf{f}(\mathbf{x}, t)$ affine in $\mathbf{x}$ ($\mathbf{f}(\mathbf{x}, t) = f(t)\mathbf{x}$).

Score-Based Generative Models Through SDEs

Continuous-Time Objective

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))} \left\| \mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) \right\|_2^2$$

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

$$\nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) = -\frac{1}{\sigma^2(t)}(\mathbf{x}(t) - \boldsymbol{\mu})$$

Note: Normality holds for $\mathbf{f}(\mathbf{x}, t)$ affine in $\mathbf{x}$ ($\mathbf{f}(\mathbf{x}, t) = f(t)\mathbf{x}$).

$$d\mathbf{x} = \sqrt{\frac{d[\sigma^2(t)]}{dt}} \cdot d\mathbf{w} \quad (\text{Variance Exploding SDE})$$

$$d\mathbf{x} = -\frac{1}{2}\beta(t)\mathbf{x}(t)dt + \sqrt{\beta(t)} \cdot d\mathbf{w} \quad (\text{Variance Preserving SDE})$$

Is it possible to explicitly derive $\boldsymbol{\mu}(\mathbf{x}(0), t)$ and $\sigma^2(t)$ for VE-SDE and VP-SDE?

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

Theorem

The moments of the SDE $d\mathbf{x} = f(t)\mathbf{x}dt + g(t)d\mathbf{w}$ satisfy:

$$\begin{aligned}\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t) \\ \frac{d\sigma^2(t)}{dt} &= 2f(t)\sigma^2(t) + g^2(t)\end{aligned}$$

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

Theorem

The moments of the SDE $d\mathbf{x} = f(t)\mathbf{x}dt + g(t)d\mathbf{w}$ satisfy:

$$\begin{aligned}\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t) \\ \frac{d\sigma^2(t)}{dt} &= 2f(t)\sigma^2(t) + g^2(t)\end{aligned}$$

Proof

$$\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] = \mathbb{E}[f(t)\mathbf{x}dt|\mathbf{x}(0)] + \mathbb{E}[g(t)d\mathbf{w}|\mathbf{x}(0)]$$

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

Theorem

The moments of the SDE $d\mathbf{x} = f(t)\mathbf{x}dt + g(t)d\mathbf{w}$ satisfy:

$$\begin{aligned}\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t) \\ \frac{d\sigma^2(t)}{dt} &= 2f(t)\sigma^2(t) + g^2(t)\end{aligned}$$

Proof

$$\begin{aligned}\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] &= \mathbb{E}[f(t)\mathbf{x}dt|\mathbf{x}(0)] + \mathbb{E}[g(t)d\mathbf{w}|\mathbf{x}(0)] \\ &= f(t)\mathbb{E}[\mathbf{x}|\mathbf{x}(0)]dt + g(t)\mathbb{E}[d\mathbf{w}|\mathbf{x}(0)]\end{aligned}$$

Score-Based Generative Models Through SDEs

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

Theorem

The moments of the SDE $d\mathbf{x} = f(t)\mathbf{x}dt + g(t)d\mathbf{w}$ satisfy:

$$\begin{aligned}\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t) \\ \frac{d\sigma^2(t)}{dt} &= 2f(t)\sigma^2(t) + g^2(t)\end{aligned}$$

Proof

$$\begin{aligned}\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] &= \mathbb{E}[f(t)\mathbf{x}dt|\mathbf{x}(0)] + \mathbb{E}[g(t)d\mathbf{w}|\mathbf{x}(0)] \\ &= f(t)\mathbb{E}[\mathbf{x}|\mathbf{x}(0)]dt + g(t)\mathbb{E}[d\mathbf{w}|\mathbf{x}(0)] \\ &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)dt\end{aligned}$$

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)$$

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)$$

Proof (Continued)

$$\mathbb{E} [d\mathbf{x}|\mathbf{x}(0)] = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)dt$$

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)$$

Proof (Continued)

$$\begin{aligned}\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)dt \\ \frac{d\mathbb{E}[\mathbf{x}(t)|\mathbf{x}(0)]}{dt} &= \frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)\end{aligned}$$

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)$$

Proof (Continued)

$$\begin{aligned}\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)dt \\ \frac{d\mathbb{E}[\mathbf{x}(t)|\mathbf{x}(0)]}{dt} &= \frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)\end{aligned}$$

Examples

$$\text{NCSN: } f(t) = 0 \quad \Rightarrow \quad \boldsymbol{\mu} = \mathbf{x}(0)$$

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)$$

Proof (Continued)

$$\begin{aligned}\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)dt \\ \frac{d\mathbb{E}[\mathbf{x}(t)|\mathbf{x}(0)]}{dt} &= \frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)\end{aligned}$$

Examples

$$\text{NCSN: } f(t) = 0 \quad \Rightarrow \quad \boldsymbol{\mu} = \mathbf{x}(0)$$

$$\text{DDPM: } f(t) = -\frac{1}{2}\beta(t) \quad \Rightarrow \quad \frac{d\boldsymbol{\mu}}{dt} = -\frac{1}{2}\beta(t)\boldsymbol{\mu}$$

Score-Based Generative Models Through SDEs

Theorem

$$\frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)$$

Proof (Continued)

$$\begin{aligned}\mathbb{E}[d\mathbf{x}|\mathbf{x}(0)] &= f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)dt \\ \frac{d\mathbb{E}[\mathbf{x}(t)|\mathbf{x}(0)]}{dt} &= \frac{d\boldsymbol{\mu}(\mathbf{x}(0), t)}{dt} = f(t)\boldsymbol{\mu}(\mathbf{x}(0), t)\end{aligned}$$

Examples

$$\text{NCSN: } f(t) = 0 \quad \Rightarrow \quad \boldsymbol{\mu} = \mathbf{x}(0)$$

$$\text{DDPM: } f(t) = -\frac{1}{2}\beta(t) \quad \Rightarrow \quad \frac{d\boldsymbol{\mu}}{dt} = -\frac{1}{2}\beta(t)\boldsymbol{\mu}$$

$$\boldsymbol{\mu} = \mathbf{x}(0) \exp\left(-\frac{1}{2} \int_0^t \beta(s) ds\right)$$

Score-Based Generative Models Through SDEs

Training

1. Sample $\mathbf{x}(0) \sim p_{\text{data}}(\mathbf{x})$, $t \sim U[0, 1]$.
2. Sample $\mathbf{x}(t) \sim q(\mathbf{x}(t)|\mathbf{x}(0))$.
3. Compute loss $\mathcal{L} = \|\mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0))\|_2^2$.

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

Score-Based Generative Models Through SDEs

Training

1. Sample $\mathbf{x}(0) \sim p_{\text{data}}(\mathbf{x})$, $t \sim U[0, 1]$.
2. Sample $\mathbf{x}(t) \sim q(\mathbf{x}(t)|\mathbf{x}(0))$.
3. Compute loss $\mathcal{L} = \|\mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0))\|_2^2$.

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

NCSN

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0), [\sigma^2(t) - \sigma^2(0)] \cdot \mathbf{I}\right)$$

Score-Based Generative Models Through SDEs

Training

1. Sample $\mathbf{x}(0) \sim p_{\text{data}}(\mathbf{x})$, $t \sim U[0, 1]$.
2. Sample $\mathbf{x}(t) \sim q(\mathbf{x}(t)|\mathbf{x}(0))$.
3. Compute loss $\mathcal{L} = \|\mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0))\|_2^2$.

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\boldsymbol{\mu}(\mathbf{x}(0), t), \sigma^2(t) \cdot \mathbf{I}\right)$$

NCSN

$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0), [\sigma^2(t) - \sigma^2(0)] \cdot \mathbf{I}\right)$$

DDPM


$$q(\mathbf{x}(t)|\mathbf{x}(0)) = \mathcal{N}\left(\mathbf{x}(0)e^{-\frac{1}{2} \int_0^t \beta(s)ds}, \left(1 - e^{-\int_0^t \beta(s)ds}\right) \cdot \mathbf{I}\right)$$

Here we omit the derivations of the variance.

Score-Based Generative Models Through SDEs

Sampling


1. Sample $\mathbf{x}(1) \sim \mathcal{N}(0, \mathbf{I})$.
2. Solve the reverse SDE using numerical solvers (SDEsolve).



Score-Based Generative Models Through SDEs

Sampling

1. Sample $\mathbf{x}(1) \sim \mathcal{N}(0, \mathbf{I})$.
2. Solve the reverse SDE using numerical solvers (SDEsolve).



- Discretizing the reverse SDE provides ancestral sampling.
- Discretizing the probability flow ODE yields deterministic sampling.

Outline

- 01 Diffusion and Score Matching SDEs (VP-SDE / VE-SDE)
- 02 Probability Flow ODE
- 03 Reverse SDE
- 04 Score-Based Generative Models Through SDEs
- 05 Flow Matching (FM)**

Continuous-Time Normalizing Flows

Let's return to ODE dynamics $\mathbf{x}_t = \mathbf{x}(t)$ in the interval $t \in [0, 1]$:

- $\mathbf{x}_0 \sim p_0(\mathbf{x})$, where $p_0(\mathbf{x})$ is a base distribution (mostly $\mathcal{N}(0, \mathbf{I})$);
- $\mathbf{x}_1 \sim p_1(\mathbf{x})$, where $p_1(\mathbf{x})$ is the true data distribution $p_{\text{data}}(\mathbf{x})$.

$$\frac{d\mathbf{x}_t}{dt} = \mathbf{v}(\mathbf{x}_t, t), \quad \text{with initial condition } \mathbf{x}(0) = \mathbf{x}_0.$$

The velocity $\mathbf{v}$ determines a flow ψ_t : $\mathbf{x}_t = \psi_t(\mathbf{x}_0)$.

Continuous-Time Normalizing Flows

Let's return to ODE dynamics $\mathbf{x}_t = \mathbf{x}(t)$ in the interval $t \in [0, 1]$:

- $\mathbf{x}_0 \sim p_0(\mathbf{x})$, where $p_0(\mathbf{x})$ is a base distribution (mostly $\mathcal{N}(0, \mathbf{I})$);
- $\mathbf{x}_1 \sim p_1(\mathbf{x})$, where $p_1(\mathbf{x})$ is the true data distribution $p_{\text{data}}(\mathbf{x})$.

$$\frac{d\mathbf{x}_t}{dt} = \mathbf{v}(\mathbf{x}_t, t), \quad \text{with initial condition } \mathbf{x}(0) = \mathbf{x}_0.$$

The velocity $\mathbf{v}$ determines a flow ψ_t : $\mathbf{x}_t = \psi_t(\mathbf{x}_0)$.

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x})) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\operatorname{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right)$$

Continuous-Time Normalizing Flows

Let's return to ODE dynamics $\mathbf{x}_t = \mathbf{x}(t)$ in the interval $t \in [0, 1]$:

- $\mathbf{x}_0 \sim p_0(\mathbf{x})$, where $p_0(\mathbf{x})$ is a base distribution (mostly $\mathcal{N}(0, \mathbf{I})$);
- $\mathbf{x}_1 \sim p_1(\mathbf{x})$, where $p_1(\mathbf{x})$ is the true data distribution $p_{\text{data}}(\mathbf{x})$.

$$\frac{d\mathbf{x}_t}{dt} = \mathbf{v}(\mathbf{x}_t, t), \quad \text{with initial condition } \mathbf{x}(0) = \mathbf{x}_0.$$

The velocity $\mathbf{v}$ determines a flow ψ_t : $\mathbf{x}_t = \psi_t(\mathbf{x}_0)$.

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x})) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\operatorname{tr} \left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)} \right)$$

- It's hard to solve the continuity equation directly due to the trace term.
- There's a method (the adjoint method) that solves this equation directly, but it's unstable and unscalable.

Continuous-Time Normalizing Flows

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x})) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\operatorname{tr}\left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

- Knowing the vector field $\mathbf{v}(\mathbf{x}, t)$, the KFP (or continuity) equation allows us to compute the density $p_t(\mathbf{x})$.
- Flow matching provides scalable approach to Neural ODEs.

Continuous-Time Normalizing Flows

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x})) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\operatorname{tr}\left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

- Knowing the vector field $\mathbf{v}(\mathbf{x}, t)$, the KFP (or continuity) equation allows us to compute the density $p_t(\mathbf{x})$.
- Flow matching provides scalable approach to Neural ODEs.

Flow Matching

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_\theta(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

Continuous-Time Normalizing Flows

KFP Theorem (Continuity Equation)

$$\frac{\partial p_t(\mathbf{x})}{\partial t} = -\operatorname{div}(\mathbf{v}(\mathbf{x}, t)p_t(\mathbf{x})) \Leftrightarrow \frac{d \log p_t(\mathbf{x}(t))}{dt} = -\operatorname{tr}\left(\frac{\partial \mathbf{v}(\mathbf{x}(t), t)}{\partial \mathbf{x}(t)}\right)$$

- Knowing the vector field $\mathbf{v}(\mathbf{x}, t)$, the KFP (or continuity) equation allows us to compute the density $p_t(\mathbf{x})$.
- Flow matching provides scalable approach to Neural ODEs.

Flow Matching

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_\theta(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

- Approximate the true vector field $\mathbf{v}(\mathbf{x}, t)$ using $\mathbf{v}_\theta(\mathbf{x}, t)$.
- Use the learned flow ψ_t for deterministic sampling: $\mathbf{x}_1 = \psi_1(\mathbf{x}_0)$.

Flow Matching: Training and Sampling

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_\theta(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

$$\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x}), \quad \mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$$

Flow Matching: Training and Sampling

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_\theta(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

$$\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x}), \quad \mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$$

Training

1. Sample $t \sim U[0, 1]$, $\mathbf{x}_t \sim p_t(\mathbf{x})$.
2. Compute loss $\mathcal{L} = \|\mathbf{v}(\mathbf{x}_t, t) - \mathbf{v}_\theta(\mathbf{x}_t, t)\|^2$.

Flow Matching: Training and Sampling

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_{\theta}(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

$$\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x}), \quad \mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$$

Training

1. Sample $t \sim U[0, 1]$, $\mathbf{x}_t \sim p_t(\mathbf{x})$.
2. Compute loss $\mathcal{L} = \|\mathbf{v}(\mathbf{x}_t, t) - \mathbf{v}_{\theta}(\mathbf{x}_t, t)\|^2$.

Sampling

1. Sample $\mathbf{x}_0 \sim \mathcal{N}(0, \mathbf{I})$.
2. Solve the ODE to obtain $\mathbf{x}_1$:

$$\mathbf{x}_1 = \psi_1(\mathbf{x}_0) = \text{ODESolve}_v(\mathbf{x}_0, \theta, t_0 = 0, t_1 = 1).$$

Flow Matching: Training and Sampling

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_\theta(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

$$\mathbf{x}_0 \sim p_0(\mathbf{x}) = p(\mathbf{x}), \quad \mathbf{x}_1 \sim p_1(\mathbf{x}) = p_{\text{data}}(\mathbf{x})$$

Training

1. Sample $t \sim U[0, 1]$, $\mathbf{x}_t \sim p_t(\mathbf{x})$.
2. Compute loss $\mathcal{L} = \|\mathbf{v}(\mathbf{x}_t, t) - \mathbf{v}_\theta(\mathbf{x}_t, t)\|^2$.

Sampling

1. Sample $\mathbf{x}_0 \sim \mathcal{N}(0, \mathbf{I})$.
2. Solve the ODE to obtain $\mathbf{x}_1$:


$$\mathbf{x}_1 = \psi_1(\mathbf{x}_0) = \text{ODESolve}_v(\mathbf{x}_0, \theta, t_0 = 0, t_1 = 1).$$

Problem: The true vector field $\mathbf{v}(\mathbf{x}, t)$ is **unknown**.

Flow Matching

$$\mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t(\mathbf{x})} \|\mathbf{v}(\mathbf{x}, t) - \mathbf{v}_\theta(\mathbf{x}, t)\|^2 \rightarrow \min_{\theta}$$

- There are infinitely many possible $\mathbf{v}(\mathbf{x}, t)$ between $p_0(\mathbf{x})$ and $p_1(\mathbf{x})$.
- We need to select the "best" $\mathbf{v}(\mathbf{x}, t)$ and make the objective tractable.



Summary

- Score matching (NCSN) and diffusion models (DDPM) are discretizations of SDEs (variance exploding and variance preserving).
- Every SDE admits a corresponding probability flow ODE that follows the same probability path $p_t(\mathbf{x})$, yielding deterministic invertible sampling and exact log-likelihoods via the instantaneous change of variables.
- Converting SDEs to PF-ODEs enables efficient high-order numerical solvers, reducing the number of sampling steps from 100–1000 to 20–50.
- SDEs can be reversed in time using the score function $\nabla_{\mathbf{x}} \log p_t(\mathbf{x})$; discretizing the reverse SDE gives ancestral sampling, while discretizing the PF-ODE gives deterministic sampling.
- The continuous-time score matching objective generalizes DDPM and NCSN; for affine-drift SDEs the transition kernel $q(\mathbf{x}(t)|\mathbf{x}(0))$ is Gaussian, making the training objective tractable.
- Flow matching fits the vector field $\mathbf{v}(\mathbf{x}, t)$ directly, an alternative parametrization of the same continuous-time generative dynamics.