

MIPT & YSDA · AUTUMN 2026

Deep Generative Models

Lecture 13

Roman Isachenko

Moscow Institute of Physics and Technology
Yandex School of Data Analysis

Recap of Previous Lecture

Let us choose $\mathbf{z} = (\mathbf{x}_0, \mathbf{x}_1)$. Then $p(\mathbf{z}) = p(\mathbf{x}_0, \mathbf{x}_1) = p_0(\mathbf{x}_0)p_1(\mathbf{x}_1)$.

$$p_0(\mathbf{x}|\mathbf{x}_0, \mathbf{x}_1) = \delta(\mathbf{x} - \mathbf{x}_0),$$

$$p_1(\mathbf{x}|\mathbf{x}_0, \mathbf{x}_1) = \delta(\mathbf{x} - \mathbf{x}_1).$$

Gaussian Conditional Probability Path


$$p_t(\mathbf{x}|\mathbf{x}_0, \mathbf{x}_1) = \mathcal{N}(\boldsymbol{\mu}_t(\mathbf{x}_0, \mathbf{x}_1), \boldsymbol{\sigma}_t^2(\mathbf{x}_0, \mathbf{x}_1)),$$

$$\mathbf{x}_t = \boldsymbol{\mu}_t(\mathbf{x}_0, \mathbf{x}_1) + \boldsymbol{\sigma}_t(\mathbf{x}_0, \mathbf{x}_1) \odot \boldsymbol{\epsilon}.$$

Recap of Previous Lecture

Let's consider straight conditional paths:

$$\mu_t(\mathbf{x}_0, \mathbf{x}_1) = t\mathbf{x}_1 + (1 - t)\mathbf{x}_0, \quad \sigma_t(\mathbf{x}_0, \mathbf{x}_1) = \epsilon.$$



Recap of Previous Lecture

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x}(0))} \mathbb{E}_{t \sim U[0,1]} \mathbb{E}_{q(\mathbf{x}(t)|\mathbf{x}(0))} \left\| \mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}(t)} \log q(\mathbf{x}(t)|\mathbf{x}(0)) \right\|_2^2$$

$$p_t(\mathbf{x}|\mathbf{x}_1) = q_{1-t}(\mathbf{x}|\mathbf{x}_0 = \mathbf{x}_1)$$

Variance Exploding SDE Probability Path


$$p_t(\mathbf{x}|\mathbf{x}_1) = \mathcal{N}(\mathbf{x}_1, \sigma_{1-t}^2 \mathbf{I})$$

$$\Rightarrow \quad \mathbf{v}(\mathbf{x}_t, \mathbf{x}_1, t) = -\frac{\sigma'_{1-t}}{\sigma_{1-t}} (\mathbf{x}_t - \mathbf{x}_1)$$


Recap of Previous Lecture

Variance Preserving SDE Probability Path

$$p_t(\mathbf{x}|\mathbf{x}_1) = \mathcal{N}(\alpha_{1-t}\mathbf{x}_1, (1 - \alpha_{1-t}^2)\mathbf{I})$$
$$\Rightarrow \mathbf{v}(\mathbf{x}_t, \mathbf{x}_1, t) = \frac{\alpha'_{1-t}}{1 - \alpha_{1-t}^2} \cdot (\alpha_{1-t}\mathbf{x}_t - \mathbf{x}_1)$$



Diffusion



OT

Recap of Previous Lecture

Continuous state space

- **Discrete time** $t \in \{0, 1, \dots, T\} \Rightarrow$ **DDPM / NCSN**.
- **Continuous time** $t \in [0, 1] \Rightarrow$ **Score-based SDE models**.

Potential advantages over left-to-right AR

- Parallel generation
- Bidirectional context
- Flexible generation order

Discrete state space

- **Discrete time** $t \in \{0, 1, \dots, T\}$.
- **Continuous time** $t \in [0, 1]$.

- Infilling and controllability
- Iterative refinement (with remasking or corrective updates)

Recap of Previous Lecture

Discrete Diffusion Markov Chain

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \text{Cat}(\mathbf{Q}_t \mathbf{x}_{t-1}),$$

Each $\mathbf{x}_t \in \{0, 1\}^K$ is a **one-hot vector** encoding the categorical state (it is just one token).

Transition Matrix

$$[\mathbf{Q}_t]_{ij} = q(x_t = i | x_{t-1} = j), \quad \sum_{i=1}^K [\mathbf{Q}_t]_{ij} = 1.$$

$$q(\mathbf{x}_t | \mathbf{x}_0) = \text{Cat}(\mathbf{Q}_{1:t} \mathbf{x}_0), \quad \mathbf{Q}_{1:t} = \mathbf{Q}_t \mathbf{Q}_{t-1} \cdots \mathbf{Q}_1.$$

- The choice of $\mathbf{Q}_t$ determines how information is erased and what the stationary distribution becomes.
- $\mathbf{Q}_t$ and $\mathbf{Q}_{1:t}$ should be easy to compute for each t .

Recap of Previous Lecture

Uniform vs. Absorbing Transition Matrix

Aspect	Uniform Diffusion	Absorbing Diffusion
$\mathbf{Q}_t$	$(1 - \beta_t)\mathbf{I} + \beta_t\mathbf{U}$	$(1 - \beta_t)\mathbf{I} + \beta_t\mathbf{e}_m\mathbf{1}^\top$
$\mathbf{Q}_{1:t}$	$\bar{\alpha}_t\mathbf{I} + (1 - \bar{\alpha}_t)\mathbf{U}$	$\bar{\alpha}_t\mathbf{I} + (1 - \bar{\alpha}_t)\mathbf{e}_m\mathbf{1}^\top$
$\mathbf{Q}_{1:\infty}$	$\mathbf{U}$	$\text{Cat}(\mathbf{e}_m)$
Interpretation	Random replacement	Gradual masking of tokens
Application	Image diffusion	Text diffusion $\approx$ Masked LM

Observation

Both schemes gradually destroy information, but differ in their stationary limit. Absorbing diffusion bridges diffusion and masked-language-model objectives.

Recap of Previous Lecture

Discrete conditioned reverse distribution

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \text{Cat} \left(\frac{\mathbf{Q}_t \mathbf{x}_t \odot \mathbf{Q}_{1:t-1} \mathbf{x}_0}{\mathbf{x}_t^\top \mathbf{Q}_{1:t} \mathbf{x}_0} \right).$$

ELBO term

$$\mathcal{L}_t = \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} D_{\text{KL}} \left(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) \right).$$

- The reverse process $p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \text{Cat}(\boldsymbol{\pi}_{\theta}(\mathbf{x}_t, t))$ is a learned categorical distribution.
- Minimizing $\mathcal{L}_t$ w.r.t. θ is equivalent to minimizing the cross-entropy

$$\mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} \text{H} \left(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0), p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) \right).$$

Recap of Previous Lecture

Independent Token-wise Forward Process

$$q(\mathbf{X}_t | \mathbf{X}_{t-1}) = \prod_{i=1}^m q(\mathbf{x}_t^i | \mathbf{x}_{t-1}^i) = \text{Cat}(\mathbf{Q}_t \mathbf{X}_{t-1})$$

$$p_{\theta}(\mathbf{X}_{t-1} | \mathbf{X}_t) = \prod_{i=1}^m p_{\theta}(\mathbf{x}_{t-1}^i | \mathbf{x}_t^i).$$

- All distributions defined by the forward process are factorized.
- Each factor conditions on the entire noisy sequence $\mathbf{X}_t$ — this is exactly the **masked language modeling** pattern.

Final objective: masked LM

$$\mathcal{L} = \sum_{t=1}^T \sum_{i=1}^m \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} \left[-\log p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_t) \right].$$

Outline

01 Discrete Diffusion

Absorbing Diffusion

Continuous Time Formulation

02 Course Overview

Outline

01 Discrete Diffusion

Absorbing Diffusion

Continuous Time Formulation

02 Course Overview

Outline


01 Discrete Diffusion

Absorbing Diffusion

Continuous Time Formulation

02 Course Overview

Generative Models Taxonomy



Absorbing Diffusion: Forward Process

Let's restrict to the case of absorbing transition matrix.

$$\mathbf{Q}_t = (1 - \beta_t)\mathbf{I} + \beta_t \mathbf{e}_m \mathbf{1}^\top, \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s),$$

$$\mathbf{Q}_{1:t} = \bar{\alpha}_t \mathbf{I} + (1 - \bar{\alpha}_t) \mathbf{e}_m \mathbf{1}^\top.$$

Each position is either still clean or already masked:

$$q(\mathbf{x}_t | \mathbf{x}_0) = \bar{\alpha}_t [\mathbf{x}_t = \mathbf{x}_0] + (1 - \bar{\alpha}_t) [\mathbf{x}_t = \mathbf{e}_m]$$

$$\mathbf{Q}_t = \begin{pmatrix} 1 - \beta_t & 0 & 0 \\ 0 & 1 - \beta_t & 0 \\ \beta_t & \beta_t & 1 \end{pmatrix} \Rightarrow \text{the masked state is absorbing.}$$

What happens in the conditioned reverse process $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$?

Absorbing Diffusion

Conditioned reverse distribution

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \begin{cases} [\mathbf{x}_{t-1} = \mathbf{x}_t], & \text{if } \mathbf{x}_t \neq \mathbf{e}_m, \\ \rho_t [\mathbf{x}_{t-1} = \mathbf{x}_0] + (1 - \rho_t) [\mathbf{x}_{t-1} = \mathbf{e}_m], & \text{if } \mathbf{x}_t = \mathbf{e}_m, \end{cases}$$

where

$$\rho_t = \frac{\beta_t \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t}.$$

- If $\mathbf{x}_t \neq \mathbf{e}_m$, then the token must be unchanged: $\mathbf{x}_{t-1} = \mathbf{x}_t$.
- Observing an unmasked token at time t fixes the entire history: $\mathbf{x}_{t-1} = \mathbf{x}_t = \dots = \mathbf{x}_0$.

Absorbing Diffusion

- If $\mathbf{x}_t = \mathbf{e}_m$, the previous token may be either clean or masked.
- With probability ρ_t , masking occurred exactly at step t (so $\mathbf{x}_{t-1} = \mathbf{x}_0$).
- With probability $(1 - \rho_t)$, the token was already masked earlier (so $\mathbf{x}_{t-1} = \mathbf{e}_m$).

Absorbing Diffusion

Sequence Distribution

$$q(\mathbf{X}_{t-1} | \mathbf{X}_t, \mathbf{X}_0) = \prod_{i=1}^m q(\mathbf{x}_{t-1}^i | \mathbf{x}_t^i, \mathbf{x}_0^i).$$


Each position i has two possible cases:

$$\mathbf{x}_t^i \neq \mathbf{e}_m \Rightarrow \mathbf{x}_{t-1}^i = \mathbf{x}_t^i, \quad \mathbf{x}_t^i = \mathbf{e}_m \Rightarrow \mathbf{x}_{t-1}^i \in \{\mathbf{x}_0^i, \mathbf{e}_m\}.$$

Interpretation

- The forward process produces **random partial observations** of the clean sequence.
- If a token is visible at time t , the reverse distribution is deterministic.
- Unmasked tokens yield a deterministic posterior and therefore contribute only a constant to the ELBO.
Therefore, only masked tokens contribute to the training loss.

Absorbing Diffusion



$$p_\theta(\mathbf{X}_{t-1}|\mathbf{X}_t) = \prod_{i=1}^m p_\theta(\mathbf{x}_{t-1}^i|\mathbf{X}_t).$$

Objective: sequence-level $\mathcal{L}_t$

$$\mathcal{L}_t = \mathbb{E}_{q(\mathbf{X}_t|\mathbf{X}_0)} \sum_{i=1}^m \rho_t[\mathbf{x}_t^i = \mathbf{e}_m] [-\log p_\theta(\mathbf{x}_0^i|\mathbf{X}_t)] + \text{const}$$

Algorithm: Discrete-Time Absorbing Diffusion

Training

1. Sample $\mathbf{X}_0 \sim p_{\text{data}}(\mathbf{X})$, $t \sim U\{1, \dots, T\}$.
2. Sample the corrupted sequence by independent masking:

$$\mathbf{x}_t^i = \begin{cases} \mathbf{e}_m, & \text{with prob. } 1 - \bar{\alpha}_t, \\ \mathbf{x}_0^i, & \text{with prob. } \bar{\alpha}_t, \end{cases} \quad i = 1, \dots, m.$$

3. Predict token distributions in parallel:

$$p_{\theta}(\mathbf{X}_0 | \mathbf{X}_t) = \prod_{i=1}^m p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_t).$$

4. Compute masked-CE loss:

$$\mathcal{L} = \rho_t \sum_{i=1}^m [\mathbf{x}_t^i = \mathbf{e}_m] \left[-\log p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_t) \right].$$

Algorithm: Discrete-Time Absorbing Diffusion

Sampling

1. Initialize $\mathbf{X}_T \leftarrow \mathbf{e}_m \mathbf{1}^\top$ (fully masked).
2. Iterate for $t = T, T - 1, \dots, 1$:
 1. Predict $p_\theta(\mathbf{x}_0^i | \mathbf{X}_t)$ for all positions.
 2. For each masked position ($\mathbf{x}_t^i = \mathbf{e}_m$):

$$\mathbf{x}_{t-1}^i = \begin{cases} \hat{\mathbf{x}}_0^i \sim p_\theta(\mathbf{x}_0^i | \mathbf{X}_t), & \text{with prob. } \rho_t, \\ \mathbf{e}_m, & \text{with prob. } 1 - \rho_t. \end{cases}$$
 3. For each unmasked position: $\mathbf{x}_{t-1}^i = \mathbf{x}_t^i$.
3. Return the final sequence $\mathbf{X}_0$ (fully unmasked).

Outline

01 Discrete Diffusion

Absorbing Diffusion

Continuous Time Formulation

02 Course Overview

From Discrete Time to Mask Rate

In absorbing diffusion, the forward process is

$$q(\mathbf{x}_t | \mathbf{x}_0) = \bar{\alpha}_t [\mathbf{x}_t = \mathbf{x}_0] + (1 - \bar{\alpha}_t) [\mathbf{x}_t = \mathbf{e}_m].$$

- The distribution depends on t only through the scalar

$$\lambda_t = 1 - \bar{\alpha}_t \in [0, 1].$$

- We can therefore reparameterize the corruption level by

$$t \quad \Rightarrow \quad \lambda \in [0, 1].$$

- We directly define a family of corrupted distributions indexed by a continuous mask rate λ :

$$q(\mathbf{x}_\lambda | \mathbf{x}_0) = (1 - \lambda) [\mathbf{x}_\lambda = \mathbf{x}_0] + \lambda [\mathbf{x}_\lambda = \mathbf{e}_m].$$

Discrete ELBO Revisited

Recall the per-step ELBO term for absorbing diffusion:

$$\mathcal{L}_t = \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} \sum_{i=1}^m \rho_t[\mathbf{x}_t^i = \mathbf{e}_m] \left[-\log p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_t) \right] + \text{const.}$$

Replacing the discrete index t with the continuous mask rate λ , the training objective becomes

$$\mathcal{L} = \int_0^1 w(\lambda) \mathbb{E}_{q_{\lambda}(\mathbf{x}_{\lambda}|\mathbf{x}_0)} \sum_{i=1}^m [\mathbf{x}_{\lambda}^i = \mathbf{e}_m] \left[-\log p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_{\lambda}) \right] d\lambda.$$

Interpretation

Training corresponds to optimizing a **continuous mixture of masked language modeling objectives** with different mask rates.

Algorithm: Masked Diffusion Language Model (MDLM)

Training

1. Sample $\mathbf{X}_0 \sim p_{\text{data}}(\mathbf{X})$, $\lambda \sim U[0, 1]$.
2. Sample the corrupted sequence by independent masking:

$$\mathbf{x}_{\lambda}^i = \begin{cases} \mathbf{e}_m, & \text{with prob. } \lambda, \\ \mathbf{x}_0^i, & \text{with prob. } 1 - \lambda, \end{cases} \quad i = 1, \dots, m.$$

3. Predict token distributions in parallel:

$$p_{\theta}(\mathbf{X}_0 | \mathbf{X}_{\lambda}) = \prod_{i=1}^m p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_{\lambda}).$$

4. Compute masked-CE loss:

$$\mathcal{L} = w(\lambda) \sum_{i=1}^m [\mathbf{x}_{\lambda}^i = \mathbf{e}_m] \left[-\log p_{\theta}(\mathbf{x}_0^i | \mathbf{X}_{\lambda}) \right].$$

Algorithm: Masked Diffusion Language Model (MDLM)

Sampling

1. Initialize $\mathbf{X}_{\lambda_1} \leftarrow \mathbf{e}_m \mathbf{1}^\top$ (fully masked).
2. Iterate for $\ell = 1, \dots, L - 1$ over a decreasing schedule $1 = \lambda_1 > \lambda_2 > \dots > \lambda_L = 0$:
 1. Predict $p_\theta(\mathbf{x}_0^i | \mathbf{X}_{\lambda_\ell})$ for all positions.
 2. For each masked position ($\mathbf{x}_{\lambda_\ell}^i = \mathbf{e}_m$):

$$\mathbf{x}_{\lambda_{\ell+1}}^i = \begin{cases} \hat{\mathbf{x}}_0^i \sim p_\theta(\mathbf{x}_0^i | \mathbf{X}_{\lambda_\ell}), & \text{with prob. } 1 - \frac{\lambda_{\ell+1}}{\lambda_\ell}, \\ \mathbf{e}_m, & \text{with prob. } \frac{\lambda_{\ell+1}}{\lambda_\ell}. \end{cases}$$
 3. For each unmasked position: $\mathbf{x}_{\lambda_{\ell+1}}^i = \mathbf{x}_{\lambda_\ell}^i$.
3. Return the final sequence $\mathbf{X}_{\lambda_L}$ (fully unmasked).

Algorithm: Masked Diffusion Language Model (MDLM)

Note: Once unmasked, a token stays unmasked. So for a currently masked token:

$$\mathbb{P}(\mathbf{x}_{\lambda_{\ell+1}}^i = \mathbf{e}_m \mid \mathbf{x}_{\lambda_{\ell}}^i = \mathbf{e}_m) = \frac{\lambda_{\ell+1}}{\lambda_{\ell}}.$$

Outline

01 Discrete Diffusion

Absorbing Diffusion

Continuous Time Formulation

02 Course Overview

Course Overview: Problem Statement

Goal

Learn a generative model $p_{\theta}(\mathbf{x})$ that matches the data distribution $p_{\text{data}}(\mathbf{x})$.

Three lenses on the same problem

- **Divergence minimization** (L1, L5):

$$\min_{\theta} D(p_{\text{data}} \parallel p_{\theta}) \quad (\text{KL, JS, Wasserstein, } \dots)$$


- **Likelihood-based** (L1–L4, L7–L8): maximize $\log p_{\theta}(\mathbf{x})$

(AR, NF, VAE, diffusion as hierarchical VAE).

- **Score-based** (L6–L10): learn $\nabla_{\mathbf{x}} \log p(\mathbf{x})$

(SM / NCSN, DDPM, score-based SDE).

Course Roadmap



Likelihood-based: Exact & Approximate Density (L1–L4)

Autoregressive (L1)

$$p_{\theta}(\mathbf{x}) = \prod_{i=1}^m p_{\theta}(x_i | \mathbf{x}_{1:i-1}) \quad \text{exact likelihood, sequential sampling}$$

Normalizing Flows (L2)

$$\mathbf{x} = \mathbf{f}_{\theta}(\mathbf{z}), \quad \log p_{\theta}(\mathbf{x}) = \log p(\mathbf{z}) + \log \left| \det \frac{\partial \mathbf{z}}{\partial \mathbf{x}} \right|$$

VAE / VQ-VAE (L3–L4)

$$\log p_{\theta}(\mathbf{x}) \geq \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}))$$

- Amortized inference + reparametrization trick.
- ELBO surgery $\Rightarrow$ optimal prior; VQ-VAE with discrete latents.

Implicit & Score-based Foundations (L5–L6)

GAN / WGAN (L5)

- Implicit p_θ , adversarial min-max

$$\min_{\theta} \max_{\phi} \mathbb{E}_{p_{\text{data}}} \log D_{\phi}(\mathbf{x}) + \mathbb{E}_{p(\mathbf{z})} \log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))$$

- Optimal discriminator $\Rightarrow$ JSD; WGAN uses Wasserstein distance via Kantorovich–Rubinstein duality.
- Evaluation tools: FID, Precision/Recall, CLIP score, human eval.

Score Matching / NCSN (L6)

- Learn $\mathbf{s}_{\theta}(\mathbf{x}) \approx \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})$ via denoising SM.
- Sample with (annealed) Langevin dynamics over multiple noise scales.

Diffusion Models (L7–L8)

Gaussian diffusion (L7)

- Forward $q(\mathbf{x}_t|\mathbf{x}_{t-1})$, direct sampling $q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$.
- Diffusion as hierarchical VAE; ELBO splits into reconstruction + prior KL + $\sum_t \mathcal{L}_t$.

DDPM (L8)

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{t, \mathbf{x}_0, \epsilon} \left\| \epsilon - \epsilon_{\theta}(\mathbf{x}_t, t) \right\|^2, \quad \epsilon_{\theta} \propto -\sigma_t \mathbf{S}_{\theta}$$

- Equivalent to score-based generative modeling (NCSN-style).

Guidance (L8)

$$\hat{\epsilon}_{\theta}(\mathbf{x}_t, t, \mathbf{y}) = (1 - \gamma) \epsilon_{\theta}(\mathbf{x}_t, t, \emptyset) + \gamma \epsilon_{\theta}(\mathbf{x}_t, t, \mathbf{y})$$

Continuous Dynamics & SDE (L9–L10)

Neural ODE / continuous-time NF (L9)

$$\frac{d\mathbf{x}(t)}{dt} = \mathbf{v}_{\theta}(\mathbf{x}(t), t), \quad \frac{d}{dt} \log p_t(\mathbf{x}(t)) = -\text{tr} \left(\frac{\partial \mathbf{v}_{\theta}}{\partial \mathbf{x}} \right)$$

Forward SDE (L9–L10)

$$d\mathbf{x} = \mathbf{f}(\mathbf{x}, t) dt + g(t) d\mathbf{w}$$

- KFP equation; VE / VP SDEs unify NCSN and DDPM.

Reverse SDE & Probability flow ODE (L10)

$$d\mathbf{x} = [\mathbf{f}(\mathbf{x}, t) - g^2(t) \mathbf{s}_{\theta}(\mathbf{x}, t)] dt + g(t) d\bar{\mathbf{w}}$$

$$\frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \mathbf{s}_{\theta}(\mathbf{x}, t) \quad (\text{deterministic equivalent})$$

Flow Matching (L10–L11)

Flow Matching objective (L10)

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{t, p_t(\mathbf{x})} \left\| \mathbf{v}_{\theta}(\mathbf{x}, t) - \mathbf{v}(\mathbf{x}, t) \right\|^2$$

Target $\mathbf{v}(\mathbf{x}, t)$ is the marginal vector field — intractable in general.

Conditional Flow Matching (L11)

$$\mathcal{L}_{\text{CFM}} = \mathbb{E}_{t, p(\mathbf{z}), p_t(\mathbf{x}|\mathbf{z})} \left\| \mathbf{v}_{\theta}(\mathbf{x}, t) - \mathbf{v}(\mathbf{x}, \mathbf{z}, t) \right\|^2$$

- $\nabla_{\theta} \mathcal{L}_{\text{FM}} = \nabla_{\theta} \mathcal{L}_{\text{CFM}}$ — regress against *conditional* field.
- Conditioning on endpoint $\mathbf{z} = \mathbf{x}_1$ (one-sided) or pair $(\mathbf{x}_0, \mathbf{x}_1)$ (two-sided).
- Straight Gaussian paths $\Rightarrow$ closed-form $\mathbf{v}(\mathbf{x}, \mathbf{z}, t) = \mathbf{x}_1 - \mathbf{x}_0$.

Discrete Diffusion (L12–L13)

Forward discrete process (L12)

Categorical $\mathbf{x}_t$ with transition matrices $\mathbf{Q}_t$:


- **Uniform:** $\mathbf{Q}_t = (1 - \beta_t)\mathbf{I} + \beta_t\mathbf{U}$
- **Absorbing:** $\mathbf{Q}_t = (1 - \beta_t)\mathbf{I} + \beta_t\mathbf{e}_m\mathbf{1}^\top$ (mask token)

Reverse process & Discrete ELBO (L12)

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \text{Cat}\left(\frac{\mathbf{Q}_t\mathbf{x}_t \odot \mathbf{Q}_{1:t-1}\mathbf{x}_0}{\mathbf{x}_t^\top \mathbf{Q}_{1:t}\mathbf{x}_0}\right),$$

each $\mathcal{L}_t$ reduces to a **cross-entropy** loss against $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$.

Generative Learning Trilemma



Generative Learning Trilemma

Rule of thumb

- **Likelihood + Coverage** $\Rightarrow$ **AR / NF**
exact density, *slow sampling*.
- **Likelihood + Fast Sampling** $\Rightarrow$ **VAE**
tractable bound, *blurry samples*.
- **Quality + Fast Sampling** $\Rightarrow$ **GAN**
sharp, *mode collapse, no likelihood*.
- **Quality + Coverage** $\Rightarrow$ **Diffusion**
stable training, *slow sampling*.
- **Quality + Faster Sampling** $\Rightarrow$ **FM / ODE**
fewer steps, *approx. likelihood*.
- **Discrete + Coverage** $\Rightarrow$ **Discrete diffusion**
parallel denoising, *iterative decoding*.

Summary

- In the absorbing case only masked tokens contribute to the ELBO, which reduces to a continuous mixture of MLM losses over the mask rate $\lambda \in [0, 1]$.
- MDLM sampling performs iterative parallel refinement from fully masked to fully unmasked sequences.
- No generative model dominates all axes of the trilemma — each family trades off likelihood, sample quality, mode coverage and sampling speed.
- Modern frontier (flow matching, discrete diffusion, distilled diffusion) tries to push the Pareto frontier of this trilemma.